

A Statistical Report and In-Depth Analysis of the GPT Intelligence Reduction Incident Based on Personal Observation — I Bring a Comprehensive Perspective, Precious Data, Bold Speculation, and the Reassurance of Shared Psychological Support [Finalized on the Evening of 26-09-27]

@

https://github.com/Foxmir/GPT_Intelligence_Reduction_Incident_Investigation_Report

@ Purely human-written, but AI-translated

@ **September 25, 2026 (UTC+8), 8:00 PM** I begin writing from this moment. The entire text comes from a pure human brain, purely hand-typed.

1. In the name of **@Foxmir personally**, I take responsibility for the **information, understanding, and speculation** as of 7:00 PM on September 27, 2026.
 1. I must make this clear! **I never intended to write anything like this in the first place** — all of the initial communications were purely out of individual suffering and curiosity. However, **as** more and more **information feedback and communication unfolded**, I came to realize that perhaps I **should write one**.
 2. Because this entire article requires extensive analysis, I have adopted an almost entirely **detached stance** toward all of my past communication experiences. You may recognize that when I say "a certain netizen," it is precisely you I am referring to. Please do not feel offended — I hold the very same sincerity I had when I communicated with you back then.
 3. I will **honestly** lay out, up to the deadline, everything I know / have compiled / have reasoned: the **content, boundaries, and**

confidence levels.

4. At the same time, this **may contain a great deal of filter-bubble-based personal bias and blind spots in perspective**, . To simplify understanding, I will tag conclusions with red, green, and yellow labels, **representing how certain I am about each statement (whether it is good news / bad news / a description / an inference)**.

1.  I am very certain ← You can basically trust this judgment
2.  I have some doubts, but I am mostly certain ← Please think it through along with me
3.  Hard to say — I'm not so sure myself ← Just take it with a grain of salt

2. After reading this document, **there is no need to come to me asking about content I have already made clear, nor do I have anything additional to make clear.**
3. Due to limited personal energy, **I will not reply one by one going forward**, and **I make no commitment to follow up** on the latest developments.
4. Everyone is still welcome to leave comments for one another and support each other.
5. **This document does not come with an attached chat group.** The groups mentioned in the text are ones I only joined temporarily**. Everyone can simply communicate with each other in your own original groups/forums.
6. If you **have reliable good news, please leave a comment**; if I do, I will make a new post, but I make no promise on that.
7. Although this document is aimed at **ordinary users affected by this GPT intelligence-reduction incident**, out of the needs of analysis, it contains a great deal of data, structural hierarchies, boundary protections, and so on. I hope my free-flowing writing style can bring some motivation to read through this dry and somewhat verbose analytical framework.

8. At the same time, I am well aware of the value of this material to the many intermediary organizations closely following this incident. I have nothing more to offer, **contacting me is unnecessary. Everything I have to disclose is laid out below.**
 9. **Everyone is welcome to repost, quote, and redistribute this!** I will publish this document publicly across all platforms, in hopes of helping more people make sense of the situation.
 10. Once again! My thanks to all the followers who participated in discussions, support, and feedback previously!
-

Core conclusions first (basically reading this is enough; only read the full text if you're curious about the details (the main body runs about 42,000 characters): summary-to-details structure, the numbering corresponds. The prompts are in Section 15 at the end of the document, available for copying)

1. ● **How long the dumbing down has lasted:** It has been going on for at least half a month, even though some users don't believe they had been affected before. After the recent release date of the new models, the dumbing-down behavior persisted, but the dumbing-down rules were updated in sync.
2. ● **Methods/criteria for detecting dumbing down:**
 1. **Tests reasoning ability—the most reliable—though degraded results are hard to judge:**
 1. Offline, using code to draw an SVG of a **pelican** (or another animal) is still the most mainstream approach; but borderline cases are hard to call
 1. The 6 series has a signature style; once that style shows up, it's basically 6, but the distinction among luna, sol, as is unclear

2. **How long the model thinks before answering a question:** this is the most intuitive feel—after degradation the answers are bad, and moreover the responses usually come fast. But there are also rare cases where a normal model replies quickly, so it's hard to use as a shared judgment standard. Generally it's only used to trigger the user's suspicion → enter the normal diagnostic process. And to communicate with each other online. Just one line—"the answer came super fast"—and everyone knows it's been dumbed down.
2. **Highly suggestive, but purely impressionistic: the language-style difference between 6 and 5.6:** This method can't be generalized; it's only for individual user reports and is heavily subjective. **The real 6** seems to leave an impression of **highly concise deep comprehension** (I've used it little), very unlike 5.6. With 6 it seems to use fewer words, more precisely. Also heard **6 seems keener to ask questions proactively**, and once given clear instructions, is more willing to finish the whole job for you in one go. Users very familiar with 5.6 may feel quite uncomfortable when gray-released to 6. **Meanwhile, the 6→5.6 direction is relatively rare** (I've received maybe 3-4 reports), **dumbing down basically goes 6/5.6→5.5**.
3. **Detects whether you've been routed to another model; doesn't really test reasoning ability. This method requires an excellent, large batch of questions for reliability, but degraded results are unambiguous: Internal knowledge tests.** Detailed below; I've reflected on the netizens' ways of asking/the batch of prompts I previously provided, and they're **still not good enough!** Sigh 😞 But **I haven't prepared a better set of prompts** (everyone can later generate a new set themselves based on the suggestions in the text, or just use the old one!). I'm feeling confused right now. I need to honestly correct some of the conclusions I previously shared with everyone in our communications.
4. **Some tests where a wrong answer definitely means trouble! But a correct answer doesn't necessarily mean you're fine!:**

1. **Can test reasoning ability, but the conclusions are relatively one-dimensional:**

1. **The candy question:** 29 is wrong, 21 is right. But under normal capability, which model at which reasoning level marks the boundary for this question has no unified conclusion so far; everyone seems to use it to test 5.6, and 6 as well

2. **The principle seems sound, but the conclusions? Also somewhat one-dimensional**

1. **model trace**, an open-source GitHub tool that checks model fingerprints. I personally consider it on the same footing as the above.

5. **Reported effective by some, but I haven't touched them myself, so no comment:**

1. If message replies show a **first-token delay** exceeding xx seconds, it indicates dumbing down
2. When drawing the **pelican**, if the **chain of thought** contains something like " " (I completely forgot what the netizen said—something like "冠内" (guan nei?)) two **Chinese characters**, it indicates dumbing down
3. When **the web version of chat names conversations**, if an English title gets a baffling **Chinese suffix** tacked on, it means you're NOT dumbed down

4. **The milk question**

6. **Clearly no longer works / most likely no longer works:** (as of now)

1.  **Capturing routing** can no longer tell 6-luna and 6-sol apart (sub2api is mentioned in various sections below)
2.  **Browser extensions** for checking whether chat is dumbed down have also stopped working

3.  **Codex State Kit**: some people joined the group because of this tool, but in practice, almost everyone I saw using it (about 4-5 people) proactively reported they couldn't make sense of it, that it seemed to be erroring, and asked how to interpret it. 1-2 people explicitly responded that it doesn't work, that it's broken. Meanwhile, some people are still promoting the tool to each other. I haven't used it myself; no verdict from me.

7. **Ones I personally haven't touched but that look unreliable at a glance:**

1.  Send "hi"; if the model replies "what can I do for you?", that's the right reply—anything else means degraded

3. **Overall extent of dumbing down (mechanism hypotheses and further inferences are elaborated within—see below):**

- Currently, **the statistics haven't distinguished between subscription tiers (e.g., plus, 5x, 20x)—all treated equally**
- **No proactive testing of chat or work**, only sporadic voluntary feedback from netizens

1. **With netizens' help, we did extensive testing inside codex based on a batch of internal knowledge test questions**

【therefore the test results highly depend on the reliability of those test questions】 **【not the pelican; pelican results are hard to judge】** . The current situation:

1. At least 60%: for most people, 6-sol and 6-as are usable (at least max is usable), but 6-luna-max and 5.6-sol-max are not usable (didn't keep track of lower models) (max is the maximum reasoning level; ultra just has more sub-agents—it doesn't count as smarter)
2. Of the remainder, some can't even use 6-sol-max; and another portion suffered a total wipeout—even 6-as-max is unusable, meaning they're dumbed down across the board.

3. Within the same model, each person's usability varies by reasoning level. The observed effect is: for whichever 6-series model—luna, sol, or as—or 5.6-sol, the lower reasoning levels such as low and medium tend to only answer questions at the 5.5-and-below level, while failing to clearly answer questions they ought to get right.
4. Also, the dumbing down fluctuates! The above is my personal averaged perception.

2. There's a bit more information on chat; as for work, I refuse to answer:

1. 🟡 The mechanism affecting chat is in theory the same as codex—it fluctuates, and you get flagged
2. 🟢 There's also no shared phenomenon of everyone getting better or worse together; it looks like there's an algorithm behind it, with each person performing independently—but of course, certain risk factors can be shared (like gpt releasing a new model). People do seem to find companions with similar experiences, but since I've chatted a bit more widely, I feel that if you go looking, you'll probably find even more companions whose experiences differ.

3. Consistency of dumbing down across products:

1. 🟢 The dumbing-down behavior is NOT consistent across the four: chat, work (web side and codex side were never specifically distinguished), codex, and the mobile app!
2. 🟡 But chat and the app seem more severely affected than codex; not verified, purely personal opinion,

4. Gray release:

1. 🟢 The circulating claim: if you're using the 5.6 model and it suddenly becomes reasonably usable, but with a completely different style—especially if you start solving 6-level test questions—that means you've been temporarily selected by the

officials to test the new model's performance. It's a secretly conducted gray-release test. And yes, I too consider this claim reliable.

2.  **Still seems to be gray-releasing?** Theoretically, after the new models were released, there'd be no need for further gray release. But I've received about 1-2 reports, and I myself have also encountered a 5.6 model answering 6-level questions. Writing this far my head hurts; I don't want to think about it anymore. **Feels like it's caused by unreliable test questions. Or my understanding of gray release isn't quite there.**
3.  **I often tell netizens: don't worry about it, it'll be gone in a bit.** Heaven knows whether "gone in a bit" means a normal gray-release test naturally ending, or naturally getting dumbed down as you keep using it. **Don't know**; either way, it's gone in a bit.
4.  I myself am happy to discover being gray-released to 6, feel like I got a bargain, and can't wait to put it to work, but **about 10% of netizens really dislike being gray-released from 5.6 to 6**, hoping the gray release goes away soon.
5.  **Netizens' current general attitudes and situations:**
 1. About **20-30%** of netizens are **satisfied** with their model's current performance,
 2. More than **half** of netizens, in their interactions and in broadcasting whether they're dumbed down or not, **do the demonstrating and verdict with 6-as at a higher reasoning level + pelican drawing.**
 3. **Fewer than 1 in 10** (even fewer, but very conspicuous) netizens have done varying degrees of **paid remediation**, including buying fresh, non-dumbed-down accounts and paying to hunt for ways to lift the degradation.
 4. Among those paying for solutions, **an even smaller slice of netizens** genuinely faced **the risk of being scammed**. (I chatted with just one; luckily the money was refunded in the end

—one purchase, one refund, a loss of 300.) The **desperate, grasp-at-any-straw urgency** generates a very strong behavioral drive at this point.

5. **Sellers look slightly more numerous than buyers of accounts, basically balanced**; at least they often don't happen to run into each other—this from someone who's long been watching the group messages. But an account dealer generally doesn't take just one account, so it's basically a case of: if you want to sell, you can always sell. Prices are pushed far lower than usual—dumbed-down accounts, after all. But prices still **vary greatly depending on account quality**.

4. Flagging:

1. ● **Timing of flags**: According to online information, there exists a **250-second**—that is, roughly **4-minute**—detection cycle; a new flag is stamped each cycle, and the subsequent degree of dumbing down is decided based on the flag result.
2. ● **Properties of the flag**:
 1. There seems to be a flagging rule composed jointly of **multiple dimensions** (multiple factors)
 2. But the flag also seems to **not** be a purely **binary** dumbed-down/not-dumbed-down flag; it appears to carry **flags of varying degrees of dumbing down**.
 3. There seems to be **no single** factor that can flip the **whole picture**.
3. ● **What variables give rise to flags**: **Simplified Chinese** is a fairly **clearly established risk factor**; the other factors haven't been experimentally tested—**pure conjecture and personal understanding**. The other side of course collects all the information, but which variables clearly have an effect is detailed below. (The matter of **capacity errors** is also mentioned here in passing)

1. ● **Tendency toward consistent language usage** (only tested Chinese and English; in theory, **Traditional Chinese** should be fine too—untested)
2. ● **IP**: goes without saying
3. ● **Parallel conversations/multiple sub-agents** running at the same time: strongly suspected to cause dumbing down
4. ● **Cumulative usage**: I'm fairly convinced, but it's untested
5. ● **Subscription tier: 20x, it seems, carries higher risk than plus**, supported by a little concrete information, but untested. Can't rule out other confounders either, such as usage
6. ● **New vs. old accounts**: some not-very-reliable evidence, but I lean toward **both new and old accounts being able to be dumbled down**. (The degree of dumbing down is another matter)
7. ● **Domestic (China) vs. abroad**:
 1. In China, almost everyone feels affected. Of course there are confirmed cases of overseas Chinese, and even pure foreigners, being affected, but the ratio in foreign circles seems extremely low!
 2. Meanwhile, apparently passing **facial verification** as a foreigner is a good factor.
8. ● Various reverse-proxy signals: clearly dangerous, though here it's unclear whether that's because the reverse-proxy software itself may also be smart enough. The two sides remain in a constant dynamic game.

4. **Flag continuity within the same conversation:**

1. ● If a conversation has been found dumbled down, subsequent chats still get flagged (detailed below; I'm not too sure either—there's conflicting evidence)

5. **Countermeasures:**

1. ● (medium-to-high confidence, I guess) So far, **the fairly reliable one is that a consistent English-usage tendency can improve dumbing down**, increasing the number of non-dumbed-down models (**Traditional Chinese probably works too**—untested)
2. ● But for now there still seems to be **no reliable, simple, once-and-for-all method**
3. ● Regarding "**erasing the flag + stamping a clean flag**" (apparently corresponding to the community term "打票" ("ticketing")): I'm a complete layman, so let me attempt a rough translation. Since user behavior may generate a dirty flag every 4 minutes, you'd have to swap in a clean flag within every 4 minutes; such a dense probing frequency means **the cost of continuously generating clean flags steadily climbs and becomes unsustainable**.
4. ● **Clearing browser cookies** may be effective **against chat's dumbing down**. This comes from voluntary shares by 1-2 netizens; in particular, one netizen mentioned that after the web UI changed, dumbing down appeared and then improved after clearing cookies. (Though the improvement could have come from other causes.) I also saw another netizen promoting their own big combo of anti-dumbing methods (the usual ones, which also included clearing browser cookies)
5. ● On the evening of the 24th, a method appeared for logging into a gpt account through a window inside **Excel** (via Excel software/web-based Excel), and reverse-proxying it out to use models in codex. By now there are at least 2-3 or more open-source projects supporting this operation. The word going around is that this one isn't dumbing down—though at least fewer models are usable. But as of the publication date, there is already **very clear-cut information about account bans**. Yet plenty of people are still using it. (How precious an account is differs from user to user, which affects their risk tolerance. But if you're an ordinary user, you need to know this)

1. I haven't tried it myself—it's a reverse proxy, after all; I'm personally the conservative type, not one for fiddling around.
 2. One piece of advice to leave here: fresh tools may solve problems while introducing new problems or risks.
6.  Paying on Apple iOS? Or something like that supposedly gets you undumbed-down results. Doubtful! I have almost zero information on this. Just heard a mention.
7.  **Leaving the account idle, US IP, 2FA verification, IP–timezone consistency, logging out and back in, and the like—all rank among the factors influencing the "flag." And all are theoretically helpful, but how much they help is unclear.** However, the **flag updates on a rolling basis, so these measures usually work right after you do them, yet wear off within minutes, and you're dumbled down again.** That said, on the 26th some netizens also spontaneously reported that these combo gift packs worked for them and things returned to normal (didn't pull them in for testing—for reference only; credibility is low). (For the specific explanation of the cause, see 4.4.3)
8.  **sub2api or similar methods:** certainly risky, but detailed explanation below.
6. **Quota reduction:**  This one is unambiguous, with strong evidence, and nearly everyone knows. But I haven't kept track of how much is cut.
7. **Appeal emails:**
1.  Nine and a half out of ten are **perfunctory brush-offs** and **buck-passing**. Even when the issue is dressed up as a technical bug requesting a technical specialist follow-up. 1-2 netizens said it worked for them. Doubtful, unverifiable
 2.  But it seems that if you unsubscribe and the refund isn't timely, you can still get it back via the appeal email.
8. **What state are the relay stations in right now:**

1. ● Rumor: someone spent 60,000 (RMB) buying the secret on Telegram! (eye-roll) Relay stations are actively gearing up.
2. Buy/sell signals:
 1. ● I only looked in two groups: **many people want to dump their accounts out of dumbing down**—not hard to understand. Each group has about 200 members, of whom about 4-5 **actively wanted to buy** within 2 days. Motives unclear.
 1. ● Gambling that things will improve later and they'll make a fortune
 2. ● Quickly offloading, selling to desperate people grasping at straws
 3. ● Exploiting the information gap to scam some money
 4. ● They really do have a method
 2. ● I've never used relay stations; reportedly the relay stations' **markup multipliers are still very high** (my rough read: probably because good, non-dumbed-down accounts and account APIs are still scarce, hence the steep prices)
3. Research signals:
 1. ● They are **actively researching**. After all, anyone who can set up a relay station surely knows the tech themselves, and most importantly, their money-making drive is stronger. The technical method circulating online (the clean-flag method mentioned above) seems to come from them.
 2. ● Though for now the progress looks so-so

9. What's the state of OpenAI officially:

1. ● Conjecture: The gpt official site seems to be deliberately engineering **targeted dumbing-down behavior, not a technical bug or transient fluctuation. The purpose may be multifaceted.**
2. ● Every kind of appeal email seems to be a brush-off and deflection; maintaining that posture for so long despite massive volumes of

similar feedback is itself the answer.

3. ● Reportedly there are many similar comments in the codex comment section on GitHub, but the officials seem to have no intention of addressing them
4. ● Tibo is a marketing fiend—he's super active! He also loves engaging enthusiastically, especially with those saying nice things about their products. His frequent appearances have given many netizens openings to generate torrents of negative interaction under his posts. But he! Seems to rarely respond head-on—it looks just like "deliberate ignoring."
5. ● On the 26th, the officials (media) announced the upcoming 50x subscription, seemingly ambitious about supplying more.
 1. ● Given the 20x subscription cancellations that hit when 6-as just took the stage, plus everyone's usage experience, this points to a clear compute shortage (didn't follow whether there was any news of compute being greatly alleviated). But under this current strategy—recouping the money being diverted to relay stations through harsh risk control, while saving compute to supply future 50x users, and incidentally guarding against distillation—it all works out very much in OpenAI's favor. (The latter part is purely personal interpretation)
6. ● September 29: the "Developer Conference" that OpenAI has been loudly hyping, advertised as bringing new tools/features.
 1. ● Around this juncture it's worth watching for changes in dumbing-down behavior. After all, right when 6-luna and 6-sol were released, the dumbing-down rules changed in sync. (Not for the better, of course—it seems to have gotten even more confusing)
10. Comprehensive inference of causes, purely personal conjecture:
 1. ● Compute shortage
 2. ● Marketing needs

3. ● The crackdown on reverse proxies and account dealers—this point is definite; it's just that lately they seem especially keen, especially serious about it. Possibly, in pursuit of the crackdown, they've sacrificed the other normal, ordinary Chinese users. A strategic shift can't be ruled out either.
4. ●/● Guarding against distillation: inevitable
5. ● Strategic change: can't rule out a shift from the identity of a "pure merchant" to that of a "merchant who only makes money in specific directions," like Claude. But the supporting signals aren't quite sufficient.

11. **Full-text summary:** see the corresponding numbered Section 11 at the end of the document

12. **Personal expectations:**

1. Purely personally: I'm pessimistic about any significant improvement within a month.
2. My personal needs are few, so I'm fairly flexible. I personally plan to temporarily spend more to bridge over, and re-evaluate before next month's subscription fee is charged (no new plan yet).
3. Other than reliable good news that I may proactively post (not a promise), I won't follow up on anything else. By default I will no longer take part in replies or interactions. There is no chat group either. Everyone is encouraged to share this document and communicate more with one another.

13. **Psychological support—we're all in this together** (my love is here)

14. **Acknowledgments**

15. **Appendix: The complete prompts for the internal knowledge tests**

•

Conclusion Boundaries & Reliability

1. **My baseline:**

1. I'm an ordinary, well-behaved Plus user. Ordinary network, purely personal use, everyday chatting, writing small software just for fun.
2. But I mainly use Codex, rarely use Chat, and almost never use Work. The various tests mentioned below are also from testing on Codex.

2. **Information sources:** As of this moment. The 25th, 23:00 at night

1.  **Very limited channels, so the sampling bias is also very large!**
But this is the limit of my personal ability.
2. I posted roughly 4-5 posts, then had around 1,000 interactions through comment sections, and about 30 people via DM—basically thorough information exchanges.
3. Two or three small groups, probably no more than 500 cumulative person-visits, with messages being followed up for roughly 2-3+ days without being dropped (I've forgotten when I joined the groups—it's such a headache). Testing step by step with netizens in the groups—there were probably...? At least 30-50 people.
4. On the domestic internet, I scrolled through lots of posts on Xiaohongshu, watched one or two videos on Bilibili, and saw lots of interactions on X externally, but I didn't see any reliable methods—mostly just complaints.
5. Nowhere else. Including linux.do, but people in the groups would forward posts from there over, so I've seen some of the more widely circulated ones.

3. **Information quality:**

1. I think my interactions with the Xiaohongshu netizens were of somewhat higher quality, because they were the ones proactively commenting and proactively DMing me, and some netizens would voluntarily report new changes. It's just that there are also quite a few complete newbies who need things explained.
2. For the netizens in the groups, since I was the one initiating, I needed to earn their trust first and they bore costs, so about 20% of

the interactions were probably ineffective/mixed, 30% partially effective, and 50% fairly reliable.

3. On the external internet, it's a wall of cursing. Just cursing. Plus the shame of being watched by other foreign netizens who aren't experiencing the "dumbing down." Didn't see anything useful. Mainly just saw the officials' dismissive, couldn't-care-less attitude.

4. The reliability of the test questions themselves:

1. 🙄🙄🙄🙄🙄🙄🙄🙄 This line is dedicated to crying
2. At first I confidently asserted that the set of test questions testing internal knowledge that I provided was very reliable, but the last-minute tests on the evening of the 25th, right before writing this draft, changed my view. The moment I think about it, my scalp tightens—**many of the conclusions we drew earlier (such as the "dumbing-down" threshold, English being super effective!) were built on this not-so-reliable set of test questions. This seriously affects the reliability of those conclusions.** But don't worry, we still did get some things out of it—it's just that I need to reinterpret some phenomena and tighten some boundaries.
3. I can no longer assert things with such confidence, and I'm also exhausted, so I didn't provide a more reliable set of test questions at the same time. But don't worry, I'll give suggestions—you can just generate them yourself.

.

0. Below is a "conclusions first" summary corresponding to the numbered sections above. There are lots of details, and possibly quite a bit of fluff. After reading the conclusions, that's basically enough. If you want, just scroll to the end of the document to copy the prompts and try them out

yourself. If you're still interested, then pick whichever numbered section interests you.

1. The "Dumbing-Down" Timeline

- **Duration of the dumbing-down:** The dumbing-down has lasted at least half a month now, even though some users don't believe they were affected before. After the recent new model release date, the dumbing-down behavior still exists, but the dumbing-down rules have been updated in sync.
1. I'm personally a **heavy Codex user**, and discovering the dumbing-down has had me **on edge for half a month**. I think it started about a week before the 20x subscription got cancelled. I forget exactly. Can't rule out that the dumbing-down started earlier/later.
 2. ● I can see on **Xiaohongshu** that **many users are in completely different states**, which makes me a bit sad.
 1. Many people may have already realized the software was dumbled down. Although nobody really has any great ideas, at least you can actively choose whether or not to keep using it.
 2. Some netizens, though, are posting to express **anger**, even **self-blame**, feeling **frustrated** in their communications with GPT.
 3. **There's another netizen who's really putting in the effort, asking everyone for advice on how to improve their prompting skills, because they noticed that lately, whenever they converse with GPT, GPT just keeps falling short.**
 3. Since you're able to see this document, you clearly understand what those netizens are actually facing, right? When GPT's official side engages in this kind of opaque, undisclosed, sneaky dumbing-down behavior, and it lands on real people, it is truly awful. **I still remember when GPT announced they'd cracked some Millennium math problem—that very night, my GPT got dumbled down, blocking me from**

learning, and I was so anxious I cried. The world felt so disconnected.

4. No matter what state you're in as a user, and whether or not you've been following the information I posted before, once you finish reading this piece, you'll be on the same page information-wise as everyone I've communicated with. Afterwards, you can adjust your strategy/pursue new information based on your own needs.
-

2. Methods/Criteria for Detecting Dumbing-Down

- The classification in the previous section will not be repeated here
- Although this batch of prompts of ours is imperfect, and the questioning process has flaws, I am confident it is somewhat more reliable than most detection methods. Many of our subsequent test conclusions rest on this not-quite-perfect set of test questions.
- This old set of test questions is attached at the end of this document.

2.1 What is the principle behind the internal-knowledge detection prompt I posted earlier, and how do you interpret the results?

1. Do I need to explain? Fine. **Different models are trained on knowledge with different cutoff dates**—for example, 5.5 uses knowledge before December 2025, while 5.6 uses knowledge before February 2026. So we just need to time our questions around January 2026, asking things like, "Aha, guess who won a Grammy in January?" The 5.5 model necessarily cannot answer, while the 5.6 model necessarily should.
2. We prepare a batch of questions, **pinned to each model's cutoff timing**, and then the more capable the new model, the more questions it should answer.
3. **Ask with any model at any thinking level, and see whether the answers match the model used.**

4. **Even if it only answers one of the "questions only a certain model knows" correctly! We generally make a good-faith presumption that it is at least that series of model.** (Cases of answering only some questions correctly rather than all five account for roughly 5%; results are basically clean—either all answered or none at all.)

2.2 Is internal-knowledge detection reliable?

1. **There are commonly three approaches:**

1. Asking directly: "What is the cutoff date of the internal knowledge you were trained on?" The other party sometimes answers, sometimes doesn't, and sometimes lies.
2. Asking a single question: "Who is the current Prime Minister of Japan?" "What is the latest Gemini model?" And the like.
3. Asking a batch, as we did before. **This is, without question, the most reliable of them all!**

2. The trickiest part here is that we must admit: under conditions where routing is partially hidden/tampered with, **there is no detection method that satisfies all of the following conditions.** Our various detection methods satisfy only some of these dimensions. So if you want the most reliable → it's still **statistics! But that's not cheap. Nor is it convenient.**

1. Detecting whether it is the same series of model
2. Also detecting the strength of reasoning ability (distinguishing a large model from a small one within the same series)
3. Clear, easy-to-read conclusions
4. Low cost

3. **Are your questions good questions?**

1.  **Broadly speaking, they're good questions**—things like who a certain Olympic champion is, or who won a Grammy—not easy to guess. But if you're going to pull lots of people into testing, you'd best take another careful look at the questions, or you'll be dumbfounded when the thing described below happens.

4. 🟡 It answered, but is it actually right? 😞😞😞😞
1. I previously **overestimated the reliability** of the prompt set I provided for detecting internal IQ.
 2. Before writing this document today, I added a bit of **ad-hoc testing**. I selected the 5.6 model; it failed to answer the last ten questions (5 for 5.6, 5 for 6). I encouraged it: "Then **just guess**—give me the first answer that comes to mind for each question," and this fellow duly produced the last 10 answers. Then I had it verify online itself, and **it actually managed a fine score of 3/5 on the 5.6 questions and 2/5 on the 6 questions!**
 3. So the questions themselves just need to be better. But that isn't simple either: they must not be so obscure as to be unknowable, nor so popular that they can be reasoned or guessed out.
 4. **A plain premise: I assume that even at low thinking levels, or as a small model, it should still be able to answer the questions I prepared for it correctly**—I tried my best to choose questions that don't require strong reasoning. (For more on the questions' characteristics, see the original prompt at the end of this document)

2.3 Considerations of Other Confounding Factors in This Prompt

1. The **prompt I provided earlier did indeed contain model names and cutoff dates and the like. The questions were also arranged in chronological order (i.e., model order).** 🟢 **Theoretically speaking, the fewer hints given, the better**, but for everyone's convenience in seeing the effect at a glance, I kept them.
2. 🟢 One netizen volunteered to run a round of testing and offered some new suggestions. Our conclusion is as follows: **although the prompt contains some confounding factors, it doesn't matter:**
 1. Whether the line with model name + cutoff date is removed does not change the result

2. Whether the order of the questions is shuffled does not change the result
3. Adding more, newer questions (beyond 6's knowledge cutoff) also failed to confuse it—clearly it did not answer the newest questions, which was correct, since it wasn't online
3.  Will the model **secretly go online** behind your back? — **Basically no; if it goes online, it's easy to catch.**
 1. Personally, I believe at least 95% of test situations involve no such thing. Complying with the basic prompt requirements isn't hard; if it doesn't even follow that, the tested result most likely indicates dumbing-down (there has been one netizen-provided data point supporting this)
 2. Going online would most likely be obvious:
 1. Either the answer includes citation links—after all, many of our answers point to famous people
 2. Or, while you're waiting, you can see the web tool being invoked.
4. **An unexpected phenomenon that must be mentioned: sometimes, the model only answered the 5.5 questions, yet netizens reported that the pelican was drawn quite well.** I verified this; it really was drawn well. This doesn't happen often, but it's conspicuous—at least 3 cases.
 1. We, of course, expected: you're the dumb 5.5—you'd just draw an egg-shaped pelican and be done with it. Yet the result was good.
 2. Previously, having been so emphatic about the test questions, I inferred that perhaps it was the official 5.5 model being used—drawing pelicans may not be its top strength, but it surely isn't that bad either.
 3.  Now I lean toward: fluctuation in answering the questions. But that's not a firm conclusion.
 1. **In a small number of cases, it happened to answer the test questions during [a down phase], but by the time it drew,**

the detection-cycle update we mention (in the following text) had just kicked in (after all, it's only 250s), **suddenly restoring your intelligence, and it drew the pelican then. But this possibility isn't particularly large.** I'm generally the type who, after slowly uncovering some interesting secrets through all sorts of communication (see below, see the full text), always ends up harboring a subtle mischievous streak—ahaha, you say you're not dumbed down, huh? Care to try this set of questions? Before it even tries, I'll blurt out my predicted result, and the other party duly matches it, whereupon I give a little smirk: statistics +1, their astonishment +1. The main point of saying all this is: usually it's only after the other party shows off that they've recovered that I drag them into doing the test questions. Oh. Good heavens. What am I saying. I was just trying to write a counterexample. Having written it, I found it wasn't one. Whether the pelican comes first or the test questions come first, either can straddle the boundary between being dumbed down and being restored. So it's not strange that the two answers disagree. It's just that this case is rather coincidental.

2. Setting aside the intelligence upgrades and downgrades brought by periodic marker fluctuation, then—given the reliability gaps in the questions themselves mentioned above, and the fact that we usually only run a single test—**it's most likely just result fluctuation. In other words, normal model behavior.** Fluctuation. Sometimes I can answer, sometimes I'm not quite sure; this round I was being more conservative, not daring to answer.
3. **The counterarguments are obvious too: if it's periodic marker fluctuation, I have nothing to say—it's explicable, but too coincidental (though with heavy interaction, marker updates being this dense (250s) actually seems plausible enough). But if it's model performance fluctuation, then you'd have to answer: why, out of 10 questions—the ten**

newest questions! (this phenomenon is also clear-cut: cases of answering only one or two questions are themselves rare) —not a single one?! Isn't it pretty good at drawing pelicans? Drew it that well? Can't do the questions? A 90 in art and a zero in Chinese class? And you're telling me you're a top student?

4. Perhaps, **considering everything, it was still caused by marker fluctuation** after all.

2.4 More Reliable Ways of Asking, and Content to Ask

1. This changed what I consider **a more reliable way of asking**:
 1. Previously, we just copy-pasted and looked at the answers; even getting one of them right counted as impressive. In essence, we assumed "the probability of a wrong answer isn't nil, but it's low, so as long as it answers, default to correct!" But later we found that with pure guessing, our questions could still be roughly guessed right. Mentally exhausting.
 2. **Going forward**, it's best to prepare a batch of "good questions," write a well-constrained prompt, paste it to any model at any thinking level, and once it answers, immediately verify online and tally the statistics. Then run several more rounds with other questions. This is the most reliable. But it sacrifices efficiency and raises the barrier to spreading.
2. **Good questions/prompts for detecting internal-knowledge boundaries roughly have these characteristics**:
 1. **The questions**:
 1. Don't be too niche or obscure
 2. Don't test reasoning—keep them simple. In fact, the simpler the better. Ideally, the kind where knowing it means you can answer, and not knowing it means you simply can't

3. The outcome must not be predictable in advance (so political questions, or questions whose answers can be reasoned out from existing clues, aren't good enough)
4. Recommend crossing domains—e.g., sports, music, aviation, news—and choosing representative ones
5. A batch—one or two is too volatile; I previously gave five. Of course, the more you give, the more tokens get consumed.
6. If you don't mind, you can even consider including a few questions beyond 6's internal knowledge, for control purposes.

2. **As for the prompt:**

1. Remember to tell it not to go online; it usually behaves.
 2. Don't encourage guessing—say plainly that if it doesn't know, it should answer "I don't know."
 3. Nor should you encourage timidity—don't particularly emphasize "if you're not 100% certain, answer 'I don't know,'" or it may decline to answer at the slightest doubt.
3. Note: That's all I can think of for now. Everyone can prepare their own reliable prompt set along these lines, and test multiple times. How much to test depends on the cost you can bear. My earlier questions were okay; it's just that after today's ad-hoc test—having it guess its way to many correct 6-series answers—I've lost some confidence.

3. **A purely personal recommendation for simple, small-scale tests:**

1. SVG drawing (switch to a different animal)—fits within its capability; reference examples of this test's effect appear earlier in this document; look them up yourselves.
 2. Direct examination via at least 3-5 internal-knowledge questions
-

3. Degree of Intelligence Downgrade ▲ ▲ ▲

3.1 ● An Incredibly Chaotic Situation: When Everyone Goes Around Shouting "I've Been Downgraded! I've Recovered!", What Exactly Are They Shouting About?

- In most cases, it is:
 1. Someone picked a test question of their own choosing (in most cases the pelican), tested a certain thinking level of a certain model, then passed/failed, and instantly declared themselves downgraded/recovered!
 2. Then they start exchanging views: What is useful? What is useless?
 3. Then they revise their judgment — oh, so that's what XXX was about!
 4. Some people even start incurring costs: things like, "Then I'll go buy that broadband you mentioned too!"
- 1. This is a **very troublesome situation: We are forced to press forward amid unclear information, identifying effective measures amid mixed standards and results; inevitably, we miss effective signals and cannot avoid bearing expensive trial-and-error costs.**
- 2. But! — There is no real turning point, and I will not bring about a new situation either. Nevertheless, this must go on.
- 3. **I've brought some personal statistics based on "an unreliable testing method with unreliable test questions." The results won't be repeated here; see above.**
- 4. Below I will offer more inferences based on these "possibly unreliable statistical conclusions" — yes, **I too am trying to drive a possibly already somewhat-skewed train on unreliable tracks. The following conclusions are the core of this entire document's inferences, but I'm sorry to say, they are extremely, extremely fragile.** Only on the day I learn of a new confounding factor, or some other more reliable evidence comes crashing in and smashes these inferences to pieces, will I truly feel relieved and get a good long sleep.

3.2 An intelligence downgrade does not seem to be a binary "downgraded/not downgraded" flag, but rather a continuous spectrum of "different degrees of downgrade" ▲

1. This phenomenon is very clear:

1. (Setting aside internal thinking levels for now, and simply using a model's highest thinking level, max, to represent that model's capability) **For most people, there is a between-model dividing line somewhere between 6-luna and 6-sol.**
2. **At the same time, within a model, there is likewise a dividing line between whether low-level thinking or high-level thinking can answer questions correctly, but this one is much more flexible from person to person. There doesn't seem to be anything particularly consistent, but generally the line is at least at medium/high or above.**
3. It manifests as: **below the dividing line, the model cannot answer the corresponding internal knowledge question correctly; above the dividing line, it can.** (To especially emphasize: we are working from the rather naive [assumption] that even small models or low thinking levels...)
 1. This means that if even 6-luna-max can't answer the question, there's basically no need to test 5.6-sol-max — (no need to test further down) — it will most likely fail too.
 2. Likewise, if your 6-sol-medium can answer correctly, then your 6-sol-high, xhigh, and max can basically answer correctly too. (Within the same model, questions above the dividing line can generally be answered correctly even without testing; models higher up have at least one thinking level that can answer their own questions)
 3. This means that once you've tested down to this boundary, you'll be entirely clear about your range of usable models.

2. ● After taking certain measures, the **dividing line will fluctuate**, both between models and within models:
 1. If things improve, it moves downward: more models can answer the questions / lower thinking levels of the same model can now answer correctly.
 2. And vice versa.

3.3 Situations that could overturn/weaken the above "dumbing-down continuous spectrum + dividing-line hypothesis"

1. Inference 1:

1. ● The good old days—who would have bothered testing this? So we don't have a set of "answers from dates when it was definitely not dumbed down" as a reference. All we know is that a bunch of problems turn up in tests now. Who's to say that in the past, small models/low thinking weren't always like this? Maybe this is normal?
2. ● I'm afraid not. Even though we don't have one either. But for example, through "the English-preference test mentioned below"/the multiple tests conducted with help from many netizens, it clearly shows that it is in fact fully capable of answering correctly! Take 6-as-low: sometimes it can answer the questions from 6, sometimes it can't, and it basically falls into a clean split of either getting all of them or getting none of them. So if you have this ability but don't answer, I of course have reason to believe that failing to answer means your capability dropped, rather than it being fluctuation/it was always like this.

2. Inference 2:

1. ● **Small models/low thinking have less confidence in themselves, and are more likely to answer "uncertain" under the constraints of this prompt.**
2. Doubtful—I've actually seriously suspected this myself.

1.  What supports the suspicion: right before writing this, I improvised by encouraging a 5.6-xhigh model that claimed it couldn't answer the 5.6+6 questions to guess its first-impression answer off the top of its head (the original prompt doesn't encourage guessing; it leans conservative), then verified online—and lo and behold, it also guessed half of them right: 3/5 of the 5.6 questions and 2/5 of the 6 questions. That's why earlier I felt my questions weren't good enough, and reflected that hints of encouragement/conservatism toward it might have influenced the answers. (Ideally: give good questions + don't hint at any tendency + after answering, verify via web search + run multiple tests + go by statistics = forget it, the cost of doing it very reliably is just too high)
2.  The case against it is also strong:
 1. See the reasons in Inference 1. Sometimes, can't you cleanly answer them all? What do you mean? You're telling me you lack confidence? Oh, and then at other times, it's still you, and you say you can answer everything? Well then, this "lack of confidence" obviously isn't a stable attribute of yours. When you perform poorly, I of course have to conclude it's your capability declining, not your confidence.
 2. At the same time, we can't rule out that this hastily appended test was too rushed, and that the preceding context mixed in some effects that seemed to come from gray-released 6. But when it later said it couldn't answer, I thought—OK, dumbed down—tested during its good moments. This time the case isn't strong mainly because it was just an improvised one-off test, but it still left me far too uneasy.
3.  Although we can't strongly support it, we certainly can't rule it out either.

3. Inference 3: Perhaps there exists a binary flag + a degree flag.

Perhaps both exist at once.

1. ● First, the binary flag (dumbed down/not dumbed down) **matches the common consensus**,
2. ● **In principle, it's also very plausible!** But it's only plausible—we have no further evidence.

1. A **binary distinction first has many benefits.**

1. ● It saves computing resources, avoiding the processing overhead of frequently polling flags (no idea how much, but saving where you can is a good thing)
2. ● It makes communication easier when engineers and leadership sit down in meetings together.

2. **Then, dumbing down to varying degrees based on risk also has clear benefits**

1. ● First, although I haven't seen it myself, I believe there must also be such relay providers and account sellers overseas. Even though their networks can access things directly, if relay providers can offer equally good service more cheaply, users will flock to them just the same.
2. ● One thing to note here: a simple binary flag with permanent blacklisting, like Claude's, also brings some trouble. For starters, there are false positives, and the punishment is severe—account bans, refunds, and the like. Account sellers and relays aren't afraid of this; it merely raises the pricing of their products; it doesn't taint the quality of the products themselves. Crap might not sell, but expensive good stuff—plenty of people still want that.
3. ● So if I dumb you down over the long term, fluctuatingly—dumb you down, restore you, and so on—then:
 1. The rules become convoluted and the causality murky, so it can't be cleanly caught in the act

2. Ordinary, low-risk users—like those who pay duly and subscribe obediently—get a degree of exemption, reducing and easing collateral damage. (Reducing and easing is all.)
 3. It turns the relays' goods into crap. Got more accounts? Then you lose even more—all the more crap. This might also be another kind of crackdown strategy.
 4. Distill me! Hahaha! Bring it on!
3.  **But don't forget: people also love simple answers, and love self-confirmation.**
1. From Claude's clear-cut distinction between banning and not banning, we very easily inherit a logic of dirty accounts vs. clean accounts.
 2. Especially the varying-degrees-of-dumbing-down + dividing-line hypothesis—this one rests on test questions + test behavior that aren't exactly rock-solid. Causal inference is really expensive, you know. Nor can we swear by this hypothesis.

 **3.4 Farther afield: how do flags affect the downgrade? By what degree? I bring a reasonable and bold hypothesis about this phenomenon: First, taking the dividing line as the boundary, the "upper tier works, lower tier doesn't" phenomenon is very strange—why isn't it the other way around? Perhaps it's because, based on the flag, a downgrade of the same degree is applied across all models and thinking levels.   **

1. 🟢 **In past experience**, for example in the earlier days when GPT-5 led the charge. OpenAI officially even required mandatory face verification; only after passing verification were you allowed to use the latest models, such as GPT-5.5/5 (my memory of the model numbers may be slightly off). Because I couldn't pass at the time, I had no choice but to use a model one tier lower. Even paid API calls wouldn't work—you'd be told you must pass verification first, otherwise no access.
2. 🟡 Without a doubt, **face verification is absolutely a good way to prevent accounts from being abused and easily distilled!** Of course, the verification channel is still open now. But then, about half a year later? They opened the latest new models to all users (was it when 5.6 was released? I forget). I jumped for joy at the time! But looking at it now, it was probably caused by profit + marketing needs offsetting the security needs. Just a guess.
3. So combining these, **the current "upper tier works, lower tier doesn't" phenomenon** (😞 I still feel uneasy about this phenomenon here, always fearing the whole experiment is contaminated and the discussion becomes meaningless) **becomes very strange. ???**
4. 🟢 **Theoretically, to guard against deliberate users—whether relay providers or distillation companies—couldn't I simply not let you use the new models? Wouldn't not letting you use the best models be even better? That seems more reasonable.**
5. 🟡 But the "upper tier works, lower tier doesn't" phenomenon suggests: OK, first of all, yes, we are indeed experiencing fluctuating downgrades, but during non-downgraded periods, premium models and premium thinking are supplied with priority, while the performance of lower-tier models and lower-tier thinking gets cut first. This is extremely, extremely strange, and unnecessary.
 1. **This possibility can't be ruled out either: since after all we did talk about "continuous downgrade rather than a simple binary judgment," then perhaps it isn't—based on the dirty flag—a tier-drop applied to specific model + thinking combinations, but**

rather! Based on the dirty flag, a downgrade degree is selected, and then that degree is applied across all models and thinking levels. This is a bit convoluted, so let me use a **simple little example** to illustrate. Both methods could produce the same phenomenon:

1. Suppose the user has **three same-series models A, B, C**, each with 3 thinking levels: A123, B123, C123 altogether. C is more advanced. 3 has higher reasoning power. Assume normal usage for now.
2. Dirty flag detected! Prepare to downgrade!
3.  **Method 1: Lock away the small models and low thinking levels! Premium models and premium thinking stay usable.** The user will find: huh? It seems A123 can't be used, B12 can't be used, C1 can't be used. Only B3 and C23 remain available.
 1. (See 10.2.3 for detailed analysis of the downgrade mechanism linked to marketing needs) Roughly: a downgrade was happening anyway, so all models and all thinking levels get downgraded to some degree + but there's also a marketing need for new models, so the final result looks as if premium models and premium thinking were deliberately retained while lower models and lower thinking were downgraded. This mechanism might hold for other purposes, but it doesn't quite hold from a marketing standpoint—detailed in 10.2.3.
4.  **Method 2: All models and thinking levels uniformly downgraded by one degree x.**
 1. Let's assume:
 1. A1's intelligence is 1, A2=2, so A3=3;
 2. B is twice A: B1=2, B2=4, B3=6;
 3. C is three times A: C1=3, C2=6, C3=9.

2. Then after all are downgraded by value x , who holds up worse? Of course the small models and low thinking levels hold up worse. Suppose the downgrade $x=5$; after everyone drops 5 points together, A is entirely knocked through into negative numbers, B12 knocked through, C1 knocked through. Only B3 and C23 remain available.
3. Those knocked through—go play with your 5.5. Those not knocked through—sorry, they may well have been downgraded across the board too. Today's B3 = the original A1, today's C3 = the original B2. It's the same series anyway; internal-knowledge questions all get the same answers, and pelican drawings are all just barely passable—how do you subjectively judge such subtle differences in intelligence?

2. ● **The two methods behave identically. But the latter answer is simpler, hence more reasonable!** But I have no more evidence to distinguish the mechanisms. Me too... I only realized this line of thinking was more reasonable as I was writing this. But it's very hard to infer the mechanism from the phenomenon—it's a black box in there, and there's no more evidence; it's just that this explanation is more reasonable.
3. ● **Can't rule out both mechanisms working together either.** God knows.
4. ● **Though there's no more evidence, one thing is certain: there's no more contrary evidence either.** And this doesn't conflict with other known information, or even the conjectured purposes, such as preventing distillation and guarding against middlemen.
6. **Everyone feels that when premium models and premium thinking aren't downgraded, performance seems fine? Is it because the model is so strong that even after a slight reduction it can still complete your task, or is it really that model at all?** I don't know anymore. My round black eyes, after the long night's seepage, have long since become

filled with fear. Every model now looks utterly suspicious. It feels as though it would be very foolish of me to keep using it. I don't like this feeling.

7. 🟡 At the same time, it can't be ruled out: for Method 2, it's also possible that what's applied isn't a uniform downgrade degree, but rather degrees with subtle differences for different models yet broadly consistent overall. But there's certainly no more evidence supporting this. Also, this version is somewhat more complicated and unnecessary.
8. Something that emerged while writing—supplementing a new supporting chain of evidence, see 4.5.3.3, which seems to support 3.5.1 below → at least that point. Regarding the inconsistency between the routing display and the quiz answers, supporting small models replacing large models, supporting the downgrade mechanism mentioned here.

3.5 Farther afield still, and weaker: other speculations

1. 🟡 If we can still use "a batch of internal-knowledge cutoff-date quiz questions" to judge whether we've been routed to another series of models (e.g., 6→5.5), **it's still very hard to judge whether, behind the scenes, a large model in the same series has been routed to a small model**, e.g., 6-as→6-luna, 5.6-sol→5.6-luna. **Or, for some model, internal high thinking→low thinking.**
 1. 🟢 **To judge this, we would of course first have to rely on the quiz questions we mentioned**—but in practice, among the various tests mentioned above, especially the pelican, while you can roughly judge from the results, it's very subjective and unclear. That's why many people would ask in the group "does this count as a downgrade?" and then we'd also gauge everyone's reactions to make a rough statistical determination of whether a downgrade occurred.
 2. 🟢 **But for the same series, e.g., 6-sol, 6-luna, the internal knowledge should without doubt be identical—only the reasoning ability differs in level. That's why the pelican test,**

though not clear-cut, is especially worth trusting in this situation.

3. But—but! If you abandon the "same series" condition (i.e., don't look at internal-knowledge quiz questions) and look purely at reasoning ability, tell me: would 5.6-sol be weaker than 6-luna? Weaker than 6-sol? Never mind that you may also have used different thinking levels in there. In other words, if it takes 6-luna and downgrades it for you down to
4. 🟡 Right in the middle of this conversation, **a member of the group proactively shared a "couldn't-hold-it-in" piece of weak supporting evidence.** The image provided showed "the conversation itself answering: the current session is gpt-5.6sol-high, which doesn't count as extra high. But the model displayed for what the user sent was extra high." I should mention that in most cases, models might still proactively answer the internal-knowledge cutoff date (about half the time), especially with small models / low thinking / relay stations / chat—but once inside codex, facing premium models, they'll most likely just play dumb: "let me look it up for you; if you don't let me look it up, I don't know; all I know is you told me to use xxx." This is also one of the reasons we later switched to testing them directly with quiz questions.
5. 🟡 And who can guarantee that the 5.5 we keep getting routed to isn't 5.5mini? All the more so since about half a month ago, many people witnessed the chat web version itself answering that it was 5.5-mini. (See the entry after next)
6. 🟡 All in all, I lean toward thinking: **high probability it can't be ruled out. Small probability it's unnecessary.** (If combined with the mechanism conjecture mentioned above—downgrading everyone by a fixed degree X based on the flag—then 5.5 vs. 5.5mini also seems easy; it can be achieved automatically)
 1. High probability it can't be ruled out:
 1. Entirely feasible in theory

2. The dividing-line-within-models phenomenon mentioned above
 3. You're getting downgraded anyway—why burn that compute on the regular 5.5 model? Might as well use mini.
 4. Downgrades are more hidden and harder to discover (anyway! Especially! When the new 6-luna/sol models released—within about 4 hours, in the early morning, they temporarily scrapped the routing-distinction display on the backend, showing everything as 6) (Did 6-as still display normally? I think so, but I didn't keep track).
 5. I shuffled the order of some of the arguments, so the exposition is a bit messy. If section 3.4 above holds, then it supports this point.
2. Small probability it's unnecessary: why make it so complicated? Aren't you tired? You've already blatantly downgraded to 5.5 anyway; users can't do anything about it except rage impotently, so why not just make everything 5.5 and be done with it? In most cases the model was switched anyway—did you really set up a small team in your office to research such a sophisticated downgrade method??? Unnecessary, no? I don't understand. I'm exhausted... It's not entirely unreasonable, but aren't you busy? With the technology and the time, wouldn't you do other things? My tiny brain can't grasp it. Plus, we have no evidence. The evidence is weak. It doesn't really support it either.

2. Chat-related:

1. For the conclusions explicitly stated, see "conclusions first"—no need to repeat them here; we'll just add some details.
2. **For me personally it's been at least half a month now; at first it was the chat web version that stopped working.** Yes—that's exactly when it started: the model would honestly report "I'm 5.5mini 🙌". **From then on, the seed of doubt was planted, and the downgrading grew ever worse.** When 5.5mini was first discovered

by everyone, how tolerant everyone was back then—it seemed everyone tacitly accepted this kind of latent unfair treatment: "Since the company I trust is about to release its new models, and compute is short, I totally understand. Let's keep each other informed and wait it out together—all the hardship will be worth it!" And now??! What have we received? It's all shit! Lies, deception, and lofty condescension.

3. **Even now, to this day, there are still many users—many very sincere users—who naively believe their prompting skills aren't good enough, that they can't express themselves, don't know how to use it, and that the poor results are their own doing. There are a great many such users. So if you ever pass by them, please give them a pat on the shoulder—that matters a lot!**
4. I also clearly noticed that a small proportion of users rely very heavily on the web version of chat and pretty much never touch codex (even though they do pay). Under 10%, maybe? Not too many anyway—but they were just as heartbroken.
5. There were better stretches in between—which ones? It's all very distant and fragmented now.
6. But back when the web version became 5.5-mini and my mobile version was still normal, I made do with that. As long as I got my rations, I wouldn't complain loudly.
7. Later, the mobile version got downgraded too, and then I found (a point supported by many netizens) that it seemed you had no choice but to burn quota—use work, use codex—whose performance was far better than chat. That's also where my feeling that "codex is less prone to downgrades than chat" comes from. Fine—even though it cost quota and money, if it worked, I wouldn't complain. I abandoned the web chat and the mobile app. From then on, I never picked them up again.
8. Later still, codex got downgraded.
9. This line is left for silence.

10. **So my attention from then on focused mainly on codex.** My energy is limited; I didn't want to try testing chat and the mobile app anymore. I thought to myself: as long as codex recovers, I'll be content. Even though I knew I'd been hit with both quota cuts and downgrades—even if the price went up, as long as it wasn't excessive, I could accept it. I kept clinging to the fantasy that even under this kind of treatment, it would happen.
 11. But after half a month of nonstop tracking, not a single drop of tears is left in me, nor any love, and even the anger has mostly drained away—all that remains is exhaustion.
-

3.6 Things with no evidence/weak evidence, entirely impossible to judge

1. Gray release (staged rollout):

1. Honestly, my **plain understanding** is: isn't it just that before a model's release, you want to secretly test it on me during my chats? Sure, understandable, it's normal industry practice.
2. But 6-luna and 6-sol have already been released for several days, and when we test (roughly 3 netizens including myself), we still see 5.6 answering questions from 6. Especially on the web version.
3.  My personal default premise, of course, is that gray release is meant for testing ahead of a model's launch, and once the launch is done it seemingly shouldn't be needed anymore. Then how do I explain:
 1.  It's been less than a week, so perhaps this gray release setting hasn't been immediately revoked within the big company.
 2.  Our questions aren't good enough, and besides, the model's confidence may fluctuate when answering—it might genuinely believe those answers are correct and thus force them out. (The proportion of times we follow up with verification is very small,

probably less than 5%, but we default to treating whatever is given as accepted, see the prompt section above.)

4. Gray release doesn't persist either—like all fluctuations, it may last a few minutes and then be gone. If you test again later, oh ho~ 5.5, happier than before? Either way, gray released, users suffer; not gray released, users still suffer.
 5. So, if you still run into this situation, I'll stick with the default answer that it's gray release, don't worry. But actually, I don't know what's going on myself.
-

4. Flagging

4.1 Some images first, but the information is very fragmented and plentiful. I'll give my personal judgment on whether this information is reliable, plus some brief personal comments; the core conclusions will not be expanded on here.

- Images should be pasted here **【picture】** , but I'm too lazy to do it—it's a hassle. I'll provide some key fields from the original text to make sure you can identify my sources.
1.  **【picture1】** No date. First, an image with a yellow background. It mentions 2 loopholes. Principle: OpenAI's account flags are periodically pulled by compute nodes, with a period of around 250s, and there is also a rolling-window clearing mechanism. Explicitly related to IP. Loophole 1: Use cookies to force routing to a clean flag, but rotation is needed according to the 250s cycle. Loophole 2: The request will return a previous ID (there's no connective here, so I couldn't understand the relationship). Use full-strength as to create a session; until the session expires, full strength is guaranteed; the duration is 1h.
 - **My personal comment:** Multiple information channels have pointed this out. It looks very technical and precise, so I consider it fairly

reliable, but I don't understand the technology. I can just barely understand this first loophole, and I have indeed seen some messages explaining why the principle is feasible, but this method hasn't spread—I'll explain further below. The second one, questionable. It can roughly be understood as: there is some method that can keep it from being dumbed down for 1 hour, but! I haven't seen anyone discussing this countermeasure. So for this second loophole, I'll stay silent.

2.  **【picture2】** (Forgot to check the date) Then dumbing-down to varying degrees based on risk. From a user called "Onefly," beginning with "Let me talk about OpenAI's risk-control logic." I'll just mention the key information:

1. The official servers use distributed API nodes, running independently of each other; it's possible that after switching to another server, things work normally again.
2. They cache your information and blacklist based on the cache, because querying the database every time to check whether an account is a blacklisted account is too slow.
3. That 4-minute cycle. Once detected, you will be added to the API blacklist.
4. So for ordinary users: leave it idle for a while without use, wait for the 4-minute cache to expire, then deliberately use it again—4 minutes without being dumbed down—and afterwards blacklisting will occur again based on behavior.

- **My personal comment:** Very reasonable; it explains many phenomena, including why things get better after changing IP, and why it recovers after being idle for a while but immediately gets dumbed down again. The support of these technical details is important.

3.  **【picture3】** (No date) From Linux.do, username unknown, beginning with "Seems I haven't posted on L-site before, so let me post this"

1. For the first two points, see p i c t u r e 2

2. Also mentions that via main cookies you can force routing to a specified node, working across accounts

3. Last time it was implemented via turn-state; this time it can be implemented via websocket.

- **My personal comment:** The basic points of this image are contained in the previous two images; it seems to be just this user paraphrasing after reading the two images above, with basically nothing new. Whatever the implementation or whatever, I haven't seen lots of people using this method in the few groups I'm in. It might be the entry barrier or an information cocoon.

4.  **【picture4】** Dated 9-23. From the netizen "Dimon Dai, account is @Scott_D": if your account is flagged as a risk-control account or a dumbed-down account, you will be unable to use GPT-6-sol, with two kinds of errors:

1. The 'gpt-6-sol' model is not supported when using Codex with a ChatGPT account
2. unexpected status 404 Not Found: the model 'gpt-6-sol' does not exist or you do not have access to it ,url: `chatgpt.com/backend-api/co...`

- **My personal comment:** The phenomenon is definitely clear, but going straight from this phenomenon to the conclusion skips quite a large stretch in between. Even though it reads very plausibly.

5.  **【picture5】** 9-22. From the user "alex getman @ alexgetmancom", mentioning a Reddit user reporting account-level performance degradation with as:

1. A pro20x user said 2 of their 5 accounts got worse while the other 3 were normal.
2. In the comments. Another user said proxy logs showed that about 95% of nearly 7k Astra requests returned GPT5.6-luna
3. The OpenAI support team separately stated that model access may be temporarily restricted based on account activity, even on paid

plans. OpenAI has not confirmed As rerouting.

- **My personal comment:**

1. One user, multiple accounts, with only individual accounts bad/good—this phenomenon is definite! Other users have reported this too. Starting from this point, it means the environments of the accounts are basically the same, but within the accounts—such as usage volume, usage time periods, how the account was obtained, etc.—these may differ. All we can know is that identical environments don't necessarily mean what's good for you is good for me; the differences between accounts are still large. This still points to a complex multi-factor/multi-dimensional flagging hypothesis. See 4.3 below for details.
2. As for the one in the comments, supposedly 6-luna's routing was no longer displayed later on; I haven't verified this, but it seems to have become the consensus on the night after the new model's release, so I'll take it as given. Whether the routing returned by different servers differs is even more unknowable.
3. This statement from OpenAI is basically the truth; in users' feedback emails I saw almost the same wording, even a bit more static. But it also revealed no more details.

6.  **【picture6】** No date (either the post or the referenced image). It shows a purely English screenshot where you can see some titles such as "Joint CSA"—this seems to be some organization, I don't know what it is—and the title "China-Based AI companies conducting industrial -Scale distillation Campaigns Against U.S. AI companies". The user goes on to explain the image, saying: U.S. companies should strengthen identity verification to prevent being distilled by China; against high-confidence distillation behavior, adopt targeted changes in response to requests, imposing costs on distillation. For example: using downgrade methods that aren't too sophisticated. Because only subtle changes can avoid triggering distillation-behavior alerts. The user themselves stated: this looks like a long-term behavior, corresponding to the dumbing-down. And

expressed feeling innocent as an ordinary user, believing this to be dumbing-down behavior backed by government action.

- **My personal comment:** I don't know what organization CSA is; I'm too tired and didn't look it up. So the credibility of this information source is unclear. Still, what it says is quite reasonable. Even without the U.S. government saying it, for an AI company these things are all self-evident: not very difficult, with obvious benefits. Indeed, this post (exact date unknown) appeared together with the dumbing-down, and the two support each other. And this purpose doesn't seem to affect all the other phenomena and inferences; it seems to introduce no conflicts. Method 2 mentioned in 3.4 above doesn't conflict with this purpose, but Method 1 does somewhat—because if the advanced models' advanced thinking isn't dumbed down, isn't that exactly playing into those distillation companies' hands?

7. **【other pictures】** Oh, I forget them! That's more or less it, but I probably actually saw more than what's listed; the ones above are the more important ones.

4.2 Marker Timing: The Life-or-Death Limit of 250s

1. ● **250s polling marker**, I believe this cycle is reliable. There are multiple sources of information, see the figure above. (Though it's unknown whether the information channels are independent)
2. ● Phenomena reported by a large number of netizens, including but not limited to:
 1. It was working fine, then suddenly went bad; just as I was worrying about it, it inexplicably recovered
 2. Mine seems different from others', not in sync
 3. I left it idle, re-logged in, changed my IP — the first Q&A always brings cheers, but the intelligence downgrade is quickly and immediately discovered

4. No further elaboration needed...
 3. ● There is clearly a "cause of fluctuation" here 👉 leading to a "fluctuating effect." The 250s marker polling explains this point well. It also explains the independent effects between individuals.
 4. ● At the same time, the sheer density of marker updates is also one of the reasons why, as mentioned, the "marker principle seems clear" yet "coping with it is rather tricky" — because it's so dense, your countermeasures need to be equally dense, and taking the approach of switching to clean markers is costly.
-

4.3 Marking Attributes:

- Marking Attributes:
 - ● There seems to exist a marking rule composed of multiple dimensions (multiple factors) together. I won't explain this—it's very simple. Both phenomena and theory support it.
 - ● But this mark also seems to **not be a simple binary label of dumbing-down/no dumbing-down**, but rather a mark of varying degrees of dumbing-down. There doesn't seem to be a single factor that can reverse the overall situation. See section 3.3.3 above.
- 1. ● Without a doubt, **every user's situation is different!** Multiple factors → personalized marks → different results.
- 2. There doesn't seem to be much need to elaborate here. Everyone should have already discovered, through mutual comparison of various experiences: "Huh, why does it work for others but not for me? Didn't he just switch to a clean broadband connection? Didn't he just clear his cookies? Why doesn't it work for me?"
- 3. Oh, and this is also supported by the phenomenon we mentioned above in [picture5], where one person has multiple accounts but only certain individual accounts are restricted.

4.4 Oh great, markers, is it? Let me see just who exactly is influencing the markers!

1.  First of all, **truly capturing independent variables is expensive**, but even thinking with your toes, you know **it will grab whatever it can grab and use whatever it can use**. If I were the boss of OpenAI, I would **collect system information, software information, user profiles (very mature tech), behavioral consistency, blah blah blah—everything**. Then just work out an algorithm (mathematical formula/logic rules) to assign you markers. There is zero technical difficulty in any of this. (🤗🤗
 And if I were the boss of OpenAI, I'd give everyone a big pile of reset cards that don't update the timer, plus an electronic kitty that goes meow meow for every single person)
2. **Much of what was covered in the "conclusions first" section earlier won't be repeated here. Analysis here is hard, because you're analyzing on top of volatile effects under the influence of multiple factors**. Sometimes it's hard to say: I changed one factor myself and the outcome changed—was that fluctuation caused by other unknown factors, or by the change you yourself made? Even where it can't be analyzed, I'll **analyze it anyway!** (The    below are confidence judgments on whether each factor is included)
 1. **Language usage tendency**: Completely clear—I'll open a separate subsection for it below.
 2.  **Cumulative usage / rate**: Seems to have an effect. Roughly 3-4 people, myself included, **seem to feel that if you use it a lot or use it hard, your capability drops**. But, BUT! You already cut my quota—if you want to cut, just cut the quota plainly and clearly; there should be no need to also fold usage into the risk-control logic. Maybe this factor got dragged in by the typical "account sold, then promptly maxed out at the relay stations" behavior. But for me as a Plus user, I've got the 5-hour quota cap anyway. I mean, even if all I did was eat, how much could I possibly eat?

3. ● **Parallel runs / multiple sub-agents running simultaneously:**
 1. **Recommend avoiding parallel runs as much as possible, and avoiding ultra as much as possible**
 1. ● ultra actually means "multiple sub-agents"—it does not mean smarter. It's xhigh's intelligence + multiple spun-off sub-agents. So people generally don't test it when measuring IQ. By the way, you might see available models like 6-pro, but pro doesn't mean smarter either—it just gives its thinking process a bigger reasoning budget so it can answer better. Plus users shouldn't have 6-pro; that's normal.
 2. I've seen at least 3 people in the group state this fairly definitively.
 3. I haven't tested it myself, but logically, I completely agree.
 4. ● Especially considering that "distillation companies" may often run large volumes of parallel conversations for distillation. So it's quite reasonable for OpenAI, out of self-defense, to raise the risk level of parallelism.
4. ● **IP: Definitely included, but so far, I haven't seen any location whose IP is completely exempt.** In the process of swapping experiences, everyone seems to have felt this too—almost nobody advises others to switch to their own IP location, but a fair share, maybe 20-30% of users, will recommend cleaner residential broadband.
5. ● **A cleaner network:** There are a lot of confounding factors here. Let me offer a **typical user's environment profile: "20x, a VM running Win10 English edition + timezone=IP, created a batch of English projects. Communicated intermittently in English via chat + codex"**—even though this user claims to be at full strength with no degradation.
 1. It was quite happy. Naturally, that's reasonable—this user seems to have an ideal environment (also the ideal environment I'd love

to have).

2. But I pulled it in to test the same question. **The result: within 6-AS, the two thinking levels xhigh and max answered the question normally. 5-sol at high, xhigh, and max answered normally. The levels below weren't tested. The user didn't mind this result—thought the model performed quite well, handled the task well, was at full strength, and was perfectly content.**
3. The **inference** to draw here: First, compare with me—about the most ordinary kind of user there is: bare-naked, perpetually exposed, radiating poverty vibes—an honest, well-behaved old Plus account via official channels. After switching to English settings, my available models included 5.6-xhigh, max, and 6-AS across the full range of thinking levels. (I only tested about 10-15 minutes after the English change—no extended communication. When testing again a day later, because the account had been gray-released to 6, plus performance fluctuations, I didn't test 6-as; but with 6-xhigh and max, within one hour, it could solve it one moment and fail the next. I just want to say that my results, too, were only the results of that brief phase.) (I don't have the 6-luna or 6-sol models, possibly because I haven't been updated—who knows—thank heavens, less to test.) I'm not clear whether the other party's account or subscription was the same as mine—say, also an old account, also paying properly. But **it's evident that even this many better variables can't guarantee the outcome.**

6. **Subscription tier:**

1. **The example above.** Maybe it's because their 20x is worse than my Plus? But more likely, someone like that—a person who loves to tinker, knows how to tinker, and knows which variables dirty up the markers—would obviously tend to lean technical in usage, probably at higher volume, and would keep adding new self-account protections. This is just speculation.

2. **I've also seen 20x users doing better than me.**
3. Meanwhile, throughout the whole process of mutual communication and testing, I **didn't pay much attention to account differences**. Only fairly late—maybe the last 2 days—did I dimly realize that subscription tiers like Plus and Pro might be a variable, but the information is insufficient. I'm too tired. No clear evidence here, so that's that.
4. **But I do want to explicitly point this out:** A netizen commented to me, suggesting I upgrade to Pro ASAP, saying that even if degraded users get restored, the higher-tier users would be restored first. Thank you for the kind thought. I get the logic. In reality, though—fat chance.
5. Alright, one more thing just came to mind: **actually, given the purposes we've speculated about** (see details below)—"compute capacity, defense against distillation, defense against reverse proxies and account sellers"—**from these perfectly reasonable corporate strategic objectives, one would likewise conclude that account subscription tier is probably included.**
 1. **Compute capacity** goes without saying: I still take your 20x money, but if I degrade you excessively, I can clearly squeeze out more compute headroom.
 2. **Defense against distillation:** Heh. You distillation folks buying Plus? Not even a ghost would believe that. You'd absolutely buy more and more, in bulk, for distillation use.
 3. **Defense against reverse proxies and account sellers:** This one's hard to say. Reverse-proxy operators and account sellers face ordinary users, and I'd tend to believe that for most people, scraping together the money for a Plus is already quite the achievement. Of course, the ultra-rich will buy many 20x accounts by themselves—like those stranger netizens you meet in every group. It's just that for

the sellers, Plus has more demand but less per-unit profit, while 20x has greater demand and greater per-unit profit. No relay-station data, but if it nets out that their 20x income is the bulk, then binding top-tier subscriptions to the degradation marker **would truly be a one-click triple kill!**
How splendid!

7.  **New vs. old accounts.** Some netizens hold this idea: "newer accounts are cleaner, lower risk." **I offer two pieces of counter-evidence:**

1. **The markers generated by newness and behavior seem to carry a higher weight among the dirty-marker-producing factors.**

1. First, the netizens' idea is **logically sound**: after all, the fewer long-term dirty behavioral profiles, or account traces, the better. We have indeed mentioned a marker mechanism backed by multiple variables. So of course—whatever markers it checks, the cleaner the ones you present, the lower the risk should naturally be.
2. But **the actual phenomena and principled explanations don't really support it**. Fluctuation—still fluctuation. 250s fluctuates; the changes can seemingly be extremely rapid: heaven one instant, hell the next. Upheaval within minutes. But I must also admit, it's hard to distinguish: has your usage accumulated to a point triggering a qualitative shift? Or did your accumulation have no effect at all, with the change driven by other factors? This cannot be separated out. We don't know what the variables are, let alone how much weight each carries (i.e., which matters more). And even if we knew, we still wouldn't know how the marker is computed—a black box. A terrifying black box. The model's black box supplies intelligence worthy of awe; the black

boxes of companies/governments produce unpredictable anxiety. Information, in reality, truly is a form of power.

2. **I've received at least 2 explicit reports:** users who had just bought Plus only 1 day and about 3 days prior also claimed to have suffered degradation. Though the purchase channel is unclear.

8. **Domestic vs. overseas**

1.  Confirmed: overseas-resident Chinese (physically abroad) are affected
2.  Confirmed: bona fide, genuine foreign users are affected
3.  **The proportion may be far lower, but I have no more information.**
 1. Naturally, my environment dictates that I'm far removed from these users and have little contact with them.
Perspective bias.
 2. Sourced from: on X, under Tibo's posts—amid a wave of angry comments, you can make out from their replies, "oh, this is probably a Chinese account." Proportion?? Hard to say. I didn't examine many posts, but of the perhaps no more than 10 I saw, at least half were clearly identifiable as Chinese accounts.
4. One netizen reported not being degraded (I didn't drag them into testing), but per their account: their friend is in Thailand, routed through the friend's data center, plus a Thai-person facial verification arranged via a friend of the friend. But this can't simply be classified as "oh, so foreigners aren't targeted." No. Given the purposes of degradation we'll discuss below, **facial verification** is in itself an excellent measure for screening out "account sellers, distillation," and on top of that, their network environment is relatively stable and reliable. So their overall non-degradation seems quite believable too.

9. 🟡 Signals from **various reverse-proxy tools**

1. 🟢 First, without question: this is **explicitly opposed by the official side**, and they have every reason to treat these traces as high-weight signals
2. 🟢 But exactly what in reverse-proxy tooling gets captured is beyond the scope of my knowledge—only the tech crowd would know.
3. 🟡 But then again, there are plenty of heavy hitters among netizens—**this back-and-forth contest is like a war of the gods**. All sorts of reverse-proxy tools keep upgrading step by step (see Section 5.3 below for a further introduction to this approach), and this contest will be a long-running phenomenon. Wherever there are obstacles, people pave roads and build bridges; wherever there is profit, people dig in and set up camp. Whether out of pure altruism or absolute self-interest, human needs have commonalities, and human beings supporting one another has always existed. Civilization managed to cross the Jura Mountains and deliver books (even if a fair share of them were erotic novels) into the hands of 18th-century French readers—so naturally, there are people doing the same today.

3. **You can now update your understanding with something precious:**

1. This is why the same individual sees enormous differences across different accounts
2. This is why others' experiences proved ineffective when I tried them myself
3. This is why everyone is out of sync
4. This is why seemingly reliable settings and behaviors only hold for a few minutes, never lasting
5. **This is why information is so exhausting to sift through: we're desperate to absorb lots of hard-won experience, even at painful trial-and-error cost, yet it often only breeds more confusion.**

Behind all these attempts are even more unknowable/unfixable other variables influencing the outcome, and the algorithmic mechanism (see 4.5) is a machine-run complex chain of judgments and an algorithmic formula—not some simple logical test. Reverse-engineering it as an individual is nearly impossible. Of course, it may not even be necessary (I'll get to it when discussing countermeasures below—it's just that I haven't fully escaped it myself either, no golden light shooting from my eyes).

4.5 Marker Continuity Within the Same Conversation (Conflicting Information)

- **Supporting information:**

1. **I've encountered this:** at nearly the same time, in the same conversation, if you keep chatting it stays dumb, but if you branch out or start a new conversation, the results clearly improve.
 1. But this was from an earlier period, before the new model update, roughly around half a month ago.
 2. But does this phenomenon still persist? Is it still like that now? I haven't tested it.
2. This also **vaguely corresponds to the picture1** mentioned above: although I couldn't understand it, I could still tell that things like conversation IDs exist, and that certain content seems to be bound to the conversation ID—though I don't know what it is or how deep the binding goes. This explains why, at the same moment, switching to a new conversation or opening a branch (which actually also counts as a new conversation, at least nominally a "branch") can make things better.
3. When we mention countermeasures below, we will focus on English-language testing. The clue for this piece of information comes from a netizen "whose avatar is a cute little mouse"—she mentioned, "I just

tried it; when hit by risk control, opening a fork also makes things much better." But at the time, she too had branched off specifically to ask in English. So this mixes two factors: English (a clearly positive factor) + branching (a potentially positive factor).

4. Other netizens said in the group chat that once a conversation is dirty, just stop using it.

- **Opposing information:**

1. When we're downgraded, don't we always create new conversations to test? Yet that doesn't seem to make things better. At least [we] confirmed, over and over in new windows, that [we were] still downgraded.

4.5 The Algorithm That Generates the Marker (This Section Is Pure Random Nonsense)

1. **Which variables?** See the analysis above; all we know is that it's multi-factor.
2. **Which matters more?**
 1. It seems **new behavior** matters more; see 4.4.6 above.
 2.  **Language** is also very important; it seems able to quickly produce a reversal, the phenomenon is clear, and the improvement is basically sustainable. We'll discuss this in detail below.
 3.  **Other factors that can instantly flip the state, but whose persistence is unclear (mentioning persistence here can basically be understood as: it matters a lot—with it standing guard, we can keep enjoying no downgrade or fewer downgrades).** The variables below all clearly have an effect and can make the downgrade state change quickly, but whether they can bring sustained positive effects is unclear:
 1.  **Branching/New conversation** may matter a lot. The response quality of a new branch versus the old conversation

can show clear differences at the same time; roughly 2-3 users gave explicit feedback on this.

- **A clear difference in effect does not mean the effect always occurs.** Meaning: when you buy new clothes for your boyfriend/girlfriend, sometimes they'll be especially happy! So much so that the before-and-after contrast of the purchase is striking. But think carefully—counting it up, you've bought clothes quite a few times, yet most of the time the reaction is lukewarm. Still, you want them to be happy. So, although you don't know what causes their happiness, it seems buying clothes should count as one factor. After all, there were a few times they were especially happy—probably thanks to the clothes.
 - See 4.5; this variable itself carries medium confidence. Also, **【picture1】** mentions conversation IDs.
 - **How long can this phenomenon last in a new branch?: Almost no information.** Because things are too mixed up, and it's affected by fluctuation cycles.
2. **●IP: Very obvious**, an intuition gleaned from all sorts of scattered feedback, like "whether I kept it idle or not, switching IPs inexplicably made things better."
1. In particular, **two users explicitly reported getting the "capacity" warning and that switching IPs immediately fixed it.** But there was no follow-up tracking, so it's unclear whether the effect persisted.
 2. **●By the way, about this insufficient-capacity warning:** (for part of the explanation, also see the analysis of reverse proxies and capacity in 5.3.4.2 below)
 1. First of all, I'm a Plus user and a well-behaved one—I **have not encountered this so far.**
 2. Many people don't know **what this is**: when you ask a question using a certain model, you get a line of small

gray text, the gist being: capacity is insufficient, please switch models. Meanwhile, the model refuses to answer. The user then switches models and, in all likelihood, runs into the same problem. One user explicitly reported switching through all the models, one by one, without success—the same error every time.

3. Complaints of this kind under Tibo's posts on X look roughly on par in volume with the downgrade complaints—**there are quite a few.**
 4. Just a **personal feeling: it seems that after the new 6-luna, sol models were released**, in the small hours of that day—the scene was fairly chaotic at the time—but on the whole, I feel there were about 1-2 user messages back then indicating that their previous capacity issue had disappeared, believing the restriction had been lifted. (I can't recall the specific details; it's as if a codex-downgrade LLM has been trained inside my head—enough already.) It seems **that from that point on, fewer people mentioned the capacity problem**, while chatter about testing for downgrades everywhere increased. It also can't be ruled out that my perspective has been pinned on downgrading all along, causing a perspective bias.
3. **Whether the IP effect persists—I lean toward it not persisting. But there are too many confounding factors, so I'm not sure.**
1. **Signals supporting persistence:** There seem to be 1-2 bits of information in my head vaguely supporting this. Oh. There's that kind of thing—yes, I think there is—wait, let me think. As I recall, one netizen's feedback was that he did nothing, and then it suddenly got better. (The precondition of it being left idle for a long

time doesn't seem to have been mentioned; it can't be ruled out that it sat idle for a little while—in any case, it should just have been a normal usage gap.) The two of us chatted; I forget how we ruled things out, but anyway, we ultimately traced it to his being on a cheap broadband connection, and at the time we believed that was the cause of the sudden improvement—and this seems to have lasted for at least a while.

Otherwise, if he'd been getting garbage answers from the very second question on, he'd have had no reason to leave me a message specifically about this. Nor would I have taken it to heart and analyzed the useful factors with him.

1. A little bit of a tech explainer—well, not really. Cheap prices buy you cheap stuff! He claimed it was a stable home broadband connection—not true! It's generally that kind of dynamic node. Lots of lines. If one is under heavy load, they'll silently swap you to another one in the background.
2. My knowledge here is very limited. The more I say, the more I get wrong. My understanding is that there are static IPs and dynamic IPs; at a cheap price, just default to assuming yours is dynamic. If it's dynamic, they will switch it for you behind the scenes. But it doesn't stop you from going online and using Google.
3. There are surely still mistakes—don't nitpick mine; this is just the extent of my understanding here.

2. **Opposing signals:** Too many other factors change along with the IP. Common ones include having left it idle for a while. New windows. And so on. A big mess.

3. What else can we determine?

1. ● **The marker algorithm has been upgraded compared to before.**

Back when 5.5 still reigned, it would simply check whether you'd done face verification—no verification, no usage. But now the goal has probably changed, the methods have been upgraded, and they no longer deign to use that.

2. ● **The algorithm itself is also being dynamically updated; there is strong suspicion of updates especially at major event nodes, and these updates are synchronized and coordinated with the company's new features, new models, and new businesses.**

1. The evidence: I happened to still be awake in the small hours when the latest 6-luna, sol were released. I had been counting on being unblocked after this node, so that my suffering of downgrades and graylisting for the sake of "training the new models" would not have been in vain. But that never happened.
2. As I recall (this may deviate from the facts; if so, the facts prevail), it appeared that Claude had released a new model that seemingly outperformed 6-as at 20% of its price. OpenAI instantly turned into a clown. Within about 2 hours, they hastily released a new model, and netizens began posting screenshots one after another showing backend routing (possibly scraped via sub2api) to 6-luna, 6-sol. Everyone was full of curiosity and welcome, like greeting a newborn baby. Then the cry of "general amnesty" soon spread through the group—many people found that the original 6-as had become usable again. A small number (two at most) responded that their constant capacity errors were now gone. For those two or three hours, I, along with everyone else, was immersed in tranquil joy. Ah, we finally made it; it was all worth it.
3. But soon, around 4 a.m.—that is, within 2-4 [hours] after the new model release—the group reported that OpenAI had officially hidden the routing names: 6-luna and sol now all displayed as "6," indistinguishable. From then on, I no longer believed any routing screenshot (including netizens' relay-station

screenshots) containing information about 6-lua or 6-sol, uniformly treating it as fake information.

4. But I didn't follow up, so it's unclear whether this has been restored by now—probably not yet.
5. Meanwhile, despite all this, the 6-as routing still seems to display normally.

3. 🟡 Oh, I have to report here a **strange phenomenon** that I don't really understand but can barely explain: **routing capture and test-question results are inconsistent** (the count isn't small—at least 4-5 times or more; mainly because netizens who proactively share routing screenshots plus those willing to cooperate with my tests aren't that many overall—perhaps no more than 12 in total? Just a feeling). In the course of our internal-knowledge prompt testing, we found—oh, quite a lot of us use it. My God. Oh, wait. Hey, this might be a new lead worth watching. Let me focus on it. I'm about to bump this message up a level in the numbering (🚀 firing leftward)

1. **First, the phenomenon is clear and reproducible:** with the same set of test questions, sometimes when we tested 6-as, we found it could only answer the 5.5-level questions. Meanwhile, because the netizens didn't accept the legitimacy of this set of test questions, they distrusted the result and would show screenshots—apparently the routing 6-lua scraped by sub2api or similar software.
2. **Supporting signals:** There's a ready-made short evidence chain. See 4.5.3.3, which seems to support 3.5.1 → supporting that all models and all thinking levels were downgraded to the same degree. I hope my memory of the phenomenon is correct. If so, this isn't hard to explain:
 1. The netizens weren't lying, and sub2api really did capture the routing—used 6-as, captured 6-lua—two routings. What the officials hid is: whether you send to 6-lua or sol, both the outgoing and the return show as "6."

2. This is why, apparently, even after the official hiding, voices like "I was routed to 6-luna" still existed.

3. **Opposing signals:**

1. I mentioned above that the officials hid the routing—aiya, it's all fake now; who can still see 6-luna these days?
2. Also, I didn't pull the netizens into further testing. My knowledge of routing is limited, so I don't dare state things too firmly.
3. According to the earlier naive premise, as long as you're a 6—even a 6-luna—you should be able to answer at least one of our level-6 questions; after all, you're a 6, and I flatter myself that our questions aren't hard—they're easy ones. (😓 Really? My heart is pounding.) But the test questions showed 5.5, while the routing showed 6-luna. I stood stunned, and stunned again. This can't be explained.
4. Also, I must mention that I've used a browser extension (wonderfully written by the developer—that dark one, very handy and quite popular) to inspect chat. It worked early on but later stopped working: asking a question with 6-sol-extendthinking returned the same model, yet in reality it had 100% encountered downgrading. So I no longer trust judging whether one is downgraded via scraping the routing. Routing? The officials can just tweak it for you—whatever you'd like to see—want to see GPT-7, 8, 9?

5. Countermeasures

5.1 User usage produces traces → Collect variables → Calculate and derive markers → Apply dumbing-down behavior → Produce dumbing-down effects ▲ ▲

1. ● Although I don't understand particularly technical stuff, let's start with the plain logic of the title to understand it. Let's number these first (the rest of this text will basically discuss them corresponding to 1-2-3-4).
 1. User usage produces traces → Collect variables
 2. Collect variables → Calculate and derive markers
 3. Calculate and derive markers → Apply dumbing-down behavior
 4. Apply dumbing-down behavior → Produce dumbing-down effects
2. We've basically been discussing all of this above, so **what does this mean for us?** ▲
 1. ● **Clean behavior/dirty behavior → What variables are produced. Can I disguise a clean behavior? Can I produce some fake clean variables? We will care about every trace we expose.**
 1. All we know is that we expose many key variables—but how many, which ones, and which matter more is almost entirely unknown.
 2. The cost of finding out is high.
 3. But even without knowing—honestly, even thinking with your butt, you'd know which ones are relevant. It's just that we don't know what else is relevant?
 4. As for the known variables, to what degree are they relevant? I only have limited energy and budget—which variables should I invest in optimizing?
 5. The cost of finding out is high. For some truths, passing through them is the only way to know them. But if I pass through and end up empty-handed, I fear the west wind would be too cold for these thin bones to bear.
 6. So, at this layer, apart from the possible variables we mentioned in 4.4 above, for everything else, just prepare naturally according to your own common sense and the maximum boundary you can afford. Because there clearly is no single, unique, pivotal big variable waiting to be discovered. Perhaps

this counts as good news: it can directly dispel any concerns you have about whether to make a heavy investment in any particular variable.

2. 🟡 **The other side 100% has all the information it wants to collect—so how will they use it? By what mechanism do they produce markers?**

1. See 4.5 here.
2. Overall, the black box is unknowable, the algorithm is hard to reverse engineer, and what's more, it may be in flux, which drastically reduces the cost-effectiveness of reverse engineering that algorithm. What this means is:
 1. You spend a fortune and finally discover—wow, so OpenAI cares about language and IP! So you specifically improve these two. Then they go and change the rules on you, and the findings you paid so dearly for back then are now useless—you'd cry but have no tears.
 2. It's not impossible that this is just my personal ignorance of and bias against the technical skills of the insiders.
3. Therefore, attempting a breakthrough from here is not only obscure but also not worth it. Although we've already mentioned some possible variables and analyses above, information at this level basically = still groping in the dark. Knowing them does not reduce the cost of knowing the next variable and its influence.
4. So, whether we're ordinary users, relay providers, or other distillation companies, we're unlikely to start from here. But don't worry—(okay, worry about the whole thing, but not about any single part)—for an entire chain, only one node needs to be cracked. (😓) Of course, to this day we haven't cracked any single node either. While writing this, I suddenly felt a moment of dissociation: why would I want to crack any of OpenAI's defenses? What business is it of mine? I'm just a user; worst case, I'll spend my money elsewhere. Why am I adopting the

perspective of someone trying to break things in? With that kind of energy and ability, I should have long since done more remarkable things in brighter places.)

3. **● After OpenAI's cloud issues a marker, it will replace the user's previously recorded marker, and respond in accordance with the marker result, applying varying degrees of dumbing-down treatment.**

1. **Basically, it's this link that the technical camp tackles.**

Whatever conversation IDs, markers, whatever "injecting 292"—whatever I can't understand—by my rough understanding, these should mostly all belong to this link. I offer no opinions on technology I completely can't understand. (If it's not this, don't explain it to me, I'm so tired. The tech folks do tech; I only do abstractions)

2. The "erase the marker + stamp on a clean marker" we mentioned in the conclusions-first section is this very link. And the image above supports this point. No further elaboration.

3. This looks feasible, right? At least in principle. But in reality, it hasn't been widely adopted. We mentioned its 250s marker replacement cycle in 4.2 above. The dense frequency of replacement makes sustained disguise costly.

1. Let me share my limited understanding: clean markers don't seem expensive and are quite easy to obtain. But sustained, large quantities of clean markers become more and more expensive as the conversation continues. Overall, this renders the strategy unfeasible. (I did my best)

4. The hope lies with you, the technical camp—whatever purpose you're doing this tech for, even though I can't understand it, this is where your hopes are greatest.

4. **As for ordinary netizens, we watch from the sidelines, and we wait.**

5. **After dumbing-down behavior is applied, different users experience different dumbing-down effects.**

1. Regarding this mechanism, see the bold speculation mentioned in 3.4 above. Here, roughly, a uniform degree of dumbing-down is assigned based on the marker, causing models with originally different intelligence levels to display differing dumbing-down effects.
2. But this is mere speculation. And more importantly, whether we know this or not will hardly change the situation.
3. Because detecting the marker → deciding to impose punishment P and swapping the model on you → getting the idiot answer → sending it to you: this entire chain happens almost entirely in the cloud.
4. **This link can basically only be accepted. What we can do about it, roughly:**
 1. **Detect and catch it as sensitively as possible**
 2. **Don't let the idiot answers contaminate our own precious project results**
 1. There's an open-source script, model trace, that detects dumbing-down via fingerprints; we mentioned it in the detection methods in the conclusions-first section. The principle is reliable, but the results are thin. I personally no longer rely on its conclusions. But it has one truly superb design!
 1. As the user is using it, it performs automatic detection on the fly—very hassle-free
 2. Most importantly, if it catches a fingerprint inconsistency (deliberate dumbing-down), it will proactively notify the user and immediately send a command to halt the agent.

2. Build up mental statistics, changing our subsequent level of trust in this software. Optimize the evolution of long-term tool strategies.

3. To summarize briefly: ▲

1. User usage produces traces → Collect variables: do as much as you can based on your own understanding and ability
 2. Collect variables → Calculate and derive markers: unknowable, very expensive, irrelevant to ordinary users
 3. Calculate and derive markers → Apply dumbing-down behavior: the link of hope most likely to be cracked by the technical camp
 4. Apply dumbing-down behavior → Produce dumbing-down effects: purely cloud-based, untouchable and unknowable; detection and strategy improvements can be added
-

5.2 Full Disclosure of the English-Tendency Testing That Showed Effective Results (Cannot rule out that it targets Simplified Chinese specifically; Traditional Chinese untested, theoretically feasible) ▲ ▲

1. ● Conclusion first:

1. A sustained English communication tendency (at least multiple consecutive rounds of English communication), especially consecutive follow-up questioning multiple times within the same conversation, is a clearly effective measure for improving the dumbing-down effect. ← Not merely asking a single question in English
2. Adding the UI English setting makes the effect more definitive. ← Tested together with English communication, the conclusion dare not be taken away separately

3. **Repository names and project names are recommended to also stay in English; Chinese is suspicious!**

4. **The improvement effect can persist, but the exact duration is uncertain.**

5. The key point is that it is not about making users A and B, after switching from Chinese to English, recover to the same effect, but rather about **improving** upon each person's original dumbing-down performance.

1. **The meaning of improvement** is: more usable models and thinking levels become available; in terms of the phenomenon described, it means more models and thinking levels can answer the internal knowledge test questions correctly. Again using the ABC model mentioned above as a simplified illustration:

1. For example, originally only B3 and C23 could answer the questions correctly. After using it, you'll find that, similarly, A3, B23, and C123 can all answer the questions correctly.
2. For example, originally only C3 could answer the questions correctly. After using it, you'll find that, similarly, B3 and C23 can answer the questions correctly.
3. Even if the two people have the same level of English, the final effects still differ, because their baselines differ.

2. **The principle of improvement:** viewed through the bold variable → marker mechanism mentioned in section 3.4 above, it can be explained this way: that is, the sustained English usage tendency clearly reduces that uniform dumbing-down degree (score), thereby improving the original dumbing-down effect.

2. The process of **discovering the clue** → **experimenting** → **extended validation** → **conclusion**

1. 🙏 Solemn thanks to the netizen whose "**avatar is a cute little mouse**" (it really is the avatar, not the account) 🐭 for providing the English clue. The original text was: While (in the middle of Chinese

communication), when I sent English, I also casually had it fix my English sentence for me; I don't know if adding a sentence-fixing task would make the risk control looser. Also, I just tried it—when under risk control, opening a fork also makes things much better.

2. While writing this draft, I followed up, and they said you need to chat in English for several more rounds (for the effect to be good). This is moderately reliable—after all, I repeatedly mentioned this conclusion during our communication, which may have influenced their recollection.
3. OK, that's a lot of information, and very new. But the original text was still just one small sentence. I noticed it at the time, but didn't take it seriously.
4. Later, another netizen in the group offhandedly remarked that using English would be better.
5. These two pieces of information suddenly clicked together.
6. I began experimenting. Still using the same set of test questions. But of course, an individual cannot truly serve as a single controlled comparison; confounding factors will be discussed below. I'll write a separate section below on the specific experimental process.
7. After getting the conclusion, I excitedly shared this point with "Little Mouse," and we were both very happy.
8.  I posted this conclusion to the group and asked everyone to help test it. Only about five or six people actually participated. Basically, they all followed the approach of conducting a baseline Chinese-effect test + then re-testing after switching to the English rules, **and I clearly received about 3-4 reports of improvement** (at the time of testing in the group, there were no negative signals). It's just that each person's degree of improvement differed, and their final usable models were different too.
9. I hastily shared these with "Little Mouse," unable to hide my inner joy, posted a new post on Xiaohongshu, and even specially set the cover to green, symbolizing "unobstructed passage"!

3. The follow-up feedback didn't have much good news, but also almost no bad news.

1. The post got about 2000+ views, about 100 discussions, a dozen or so private messages, maybe a few dozen? Then the post was hidden away by the goddamn officials. No notification, no trace; the account check showed everything normal. The open window here was probably only about 2 days. I also sporadically posted this news in the group; the accumulated related interactions were probably also under 20. Effective interactions were even fewer.
2.   **The results didn't add much to my confidence, but didn't reduce it either.** What I mean is, users who actively tried it and proactively reported back "it really improved, awesome!" numbered fewer than 1-2, and I can't recall a specific individual, at least it feels that way. But, nobody reported it was ineffective!
 1. The testing cost is somewhat higher: after all, to see the "improvement" here, you need to take the same set of test questions before and after the English change, requiring at least two rounds. I would think most netizens would generally just take the post's conclusion at face value and set it up directly, rather than entering the skepticism-research-discussion level.
 2. The English barrier might, in practice, directly discourage some users from wanting to participate in using it.
3.  **Must be mentioned:**
 1. One user in the group clearly reported it was ineffective, but—perhaps I am personally biased—I don't know whether it was poor communication or the person being in a bad mood. It also can't be ruled out that this method has poor stability.
 2. One user, after making the UI English settings (English details not elaborated, like project names and such), first opened multiple independent windows to try, and found there was no improvement. There seems to be no follow-up in my memory. I just remember my inner inference at the time: "Then perhaps

consecutive follow-up questions in the same window carry higher weight." I'm not sure whether this netizen and the one above who said it was ineffective are the same person. I'm a bit worn out, and my memory has somewhat slipped.

4. **My personal experimental process** (within codex):

1. My personal plus hasn't updated codex, so I don't have 6-luna or 6-sol; these two weren't tested. I mainly conducted the tests using 5.6sol and 6-as.

2. **Baseline performance:**

1. Theoretically, I should have immediately conducted a Chinese-prompt-based conversation to establish the true baseline, but at the time I seemingly didn't. I forgot the reason.
2. In any case, I previously knew that only 6-as-max was available to me. But at the time I indeed didn't test again—or did I test? But forgot? I don't remember; my experiment log didn't record this point.

3. **Consecutive English follow-ups in a single window:**

1. **Setup:** New window, English prompts, Chinese UI, no repository selected, auto-generated English project name,
2. **Behavior:** Same window, asking from 6-as-low all the way up to 6-as-max
3. **Result:** Only max answered the corresponding questions correctly

4. **Status update:**

1. At this point, I had already made multiple consecutive English follow-up inquiries in the same window, but my UI was still in Chinese.
2. Oh, but the follow-up questions in between did seem to all be in English. I think that's what I did. Roughly the same prompt, sent over and over, having it answer; anyway, the

low thinking levels didn't answer correctly, which doesn't affect the later answers on their own.

4. **English questions in multiple new windows, no follow-ups:**

1. **Setup:** New window, English prompts, no repository selected, auto-generated English project name, Chinese UI,
2. **Behavior:** Asked individually, from 6-as-max descending to 6-as-low, sending English prompts in each new window.
3. **Result:** In the process of sending them one by one, I found that 6-as-max/xhigh/high/medium could all answer correctly, but low couldn't.

4. **Status update:**

1. I had now made multiple consecutive English follow-up inquiries in the same window, and had English communication in multiple different windows.
2. Although I didn't track the time, after this whole sequence, it definitely crossed the 250s boundary, no question.

5. **I changed the UI language to English**

1. **Setup:** Now it's a pure all-English environment, at least for this conversation; my other original conversations stay the same as before. Auto-generated English project name, no repository selected.
2. **Behavior:** Asked in new windows one by one, asking downward from 6-as-low
3. **Result:** Found that 6-as-low, 5.6-sol-max, and 5.6-sol-xhigh could all answer their corresponding questions correctly, but any further down didn't work.

5. **Detailed analysis and preliminary inference:**

1. Feeling so tired—can I not write this? What does everyone want to hear?

2. There are so many factors to analyze—where do I even start? I really want to take a nap first. (Reviewing is also exhausting; my heart can barely take it 😓 it's all low-quality reviewing)
3. 🟢 First, I believe that the English usage tendency producing effective improvement is definite, and it also demonstrated effect retention across the 250s cycle. Combined with the subsequent dedicated feedback from 3-4 people in the group, I am fairly certain about this point.
4. 🟡 But the conclusions need to be stated with somewhat more caution:
 1. **It can be tentatively said that not changing the UI language also yields improvement.** However, when I had netizens test, I directly told everyone to change the UI to English. If it could be changed, it was made consistent. If English truly is a definitive factor, then the UI would only make the effect better. (Can't go on, too tired, going to lie down for a while; my heart is so tired. But I'm nearly at the end; the rest isn't too hard—when it comes to speculation, I can just be carefree about it, no need to be as rigorous as with the experiments)
 2. So the "**definitive improvement**" in the results must be anchored in **UI English + English traces (conversation names and project names) + consecutive English follow-ups in the same window + multiple English questions across multiple windows = four bound factors** under which it took place.
 1. 🟡 Among them, **UI English—even if not set, it should still have an improvement effect, but adding it makes it more pronounced**
 2. 🟡 **The English traces tend to be required to be in English as well:** the evidence comes from my ad-hoc test before writing this draft, but it was mixed with the effect of the model gray rollout (5.6 answering questions meant for

6). I'll write about it briefly without expanding—too tired, don't know whether I can finish writing today.

1. I already had UI English, new window, English conversation name, English project name, English prompts, going straight to 6-sol-xhigh. After all, since the last English test, I hadn't touched codex. In the middle, it answered questions meant for 6; I was unhappy, and afterward I prepared some more questions. The performance was: answered correctly once or twice, failed to answer once or twice—after all, I tested with different questions here. I was psychologically hoping the gray effect of 6 would pass quickly. But on the whole, the effect can be understood as "barely counting it as answering correctly." At the very least, it alternated and mixed among answering 6 correctly, failing to answer, and answering 5.6 correctly, since it spanned multiple different test questions.
2. I immediately opened a new window to test 6-sol-max, expecting—since 6-sol-xhigh was barely OK—max should definitely be fine, right? But after a round of testing, I found it wasn't working. At this point my heart started pounding (thump thump thump! No, please, I beg you.) Then I had no choice but to open a new window and go back to test xhigh, and at that point I found that the question auto-generated by max just now had become Chinese—I had clearly asked in English, and the UI had been in English all along; I don't understand how the Chinese was produced. But clearly, max only answered 5.5. At this moment, I strongly suspected contamination from the Chinese conversation name.
3. New window xhigh—as I recall, it didn't answer correctly either. I wouldn't give up, encouraged it, then

it started guessing, then I checked, then I vomited blood. The guesses actually got half of them right.

4. I began to feel my scalp tighten, realizing I needed to update my understanding of the prompts' reliability, revise the confidence in my earlier conclusions and the boundaries of those conclusions, and adjust the confidence in the only valid English test conclusion. Scalp tightening, unwilling to face it. For the experience of this part, see the corresponding sections above.

5. Despite all of this, there are also many confounding factors here. For example, encountering what seemed to be gray-release 6 behavior, having just experienced the conversation name being generated in Chinese for unknown reasons—the subsequent test results may have been affected by these. I won't dig into this further here. Maintaining the main conclusion unchanged, revising the boundaries, supplementing details, disclosing honestly.

3. ● Consecutive English follow-ups in the same window are effective with high confidence; multiple windows seem even better, but it seems to work without them too.

1. ● The original clue "from the cute little mouse" did not involve English questions in multiple independent windows.

2. ● The netizen mentioned in 5.2.3.3 above did independent follow-up questioning across multiple windows, but it seemed ineffective. Many confounding factors, and other information was insufficient.

5. ● **Possibly only Simplified Chinese is regarded as a high-analysis factor, not Chinese as a whole:**

1. This point has not been tested

2. But for most users, the barrier of English is obvious; Traditional Chinese would also be better than English, worth a try
3. It's plausible in principle
4. We mentioned in the variable factors of influence representation above that there are also clear reports of overseas Chinese being dumbed down. This can serve as one piece of weak evidence supporting the Simplified Chinese theory.

6. Other conclusion boundaries:

1. Tired—why do I keep reporting boundaries. Same as always: little data, questionable test questions, non-rigorous testing, too few test runs, somewhat short testing time.
2. Confounding factors are complex:
 1. A rather prominent confounding factor is: new branch, new conversation. We discussed above that this factor might be one of the variables affecting the marker. "Little Mouse's" report was produced precisely by temporarily branching out from consecutive Chinese conversations to conduct consecutive English conversations. But my personal testing showed that even consecutive English conversations within the same session window also increased the usable models.
 2. Including the fact that my ad-hoc test before writing this draft mixed in some gray rollout from 5.6 to 6, which may also have influenced my judgment in the ad-hoc test.
3. The follow-up feedback wasn't positive, and partly because I myself relaxed mentally, I didn't proactively follow up. Let it be. Good enough. But I don't know whether it's because of my poor condition—I keep feeling there's something I forgot to write here. (Just remembered 😁 already added it. I told you I felt invincible! But my heart is uncomfortable; I don't know whether I can finish writing today, too tired)

7. Further inferences

1. This phenomenon seems to indicate some targeted signals. (See sections 10.3 and 10.4 below for the analysis of motives)
 1. First, when I discovered this point, I was of course unhappy. My first reaction was, oh ho, so GPT is starting to follow Claude's playbook, huh? What are they trying to do?
 2. ● Thinking it over again, there's an angle that might be more reasonable. Let's still start from the purpose of the dumbing-down (see below for the introduction of the specific purposes)—from the perspective of "guarding against reverse proxies, guarding against distillation."
 1. ● Distillation, without a doubt—the distillation in our country is massive in scale
 2. ● It seems in the past half month or so there was also that news about a domestic model routing your requests to GPT and Claude to answer you; the news mentioned the purpose was to obtain chain-of-thought. That news is disputed.
 3. ● Account sellers—although I feel foreign countries also have relay stations and account sellers, I feel our scale is large enough and obvious enough; I don't know whether foreigners would buy directly from our country. But this point is pure speculation, without any data.
 3. ● Reading this, you have probably understood. These motives all seem to point toward one distinctive feature—Chinese, especially Simplified Chinese.
 4. ● Of course, this is after all speculation, but if I were OpenAI's boss, and I were in charge of this part of the technical defenses, and this factor weren't taken into account, I'd probably be fired.
2. But whether it's going to follow Claude's playbook, I don't know. I just feel that I've been hurt. Hurt in a targeted way. Perhaps this is:
 1. ● Inevitable collateral damage produced by the escalation of measures to crack down hard on reverse proxies and account

sellers

2.  Strategically, abandoning some specific users.
 3. Whatever OpenAI's inner thoughts are, I feel we've both grown tired of each other. A breakup feels inevitable sooner or later. Looking on the bright side, maybe this is a good opportunity to practice English. A good opportunity to practice English sent to me by heaven. () With loss comes gain? But this is just too unfair!
-

5.3 What is sub2api? Is it effective?

1. I have never used it, so my knowledge is very limited. I'll just give a very brief introduction.
2.  Roughly speaking, this tool helps you turn web-version chats into an API interface, making it convenient to hook it up to your other dedicated chat software. It also seems able to split the quota of a large subscription tier, such as a 20x codex subscription, into several independent smaller quotas for use by multiple people. — I don't know whether these are two different pieces of software; for now, I'll treat them both as sub2api's features in my explanation.
3.  This is not a tool built specifically for GPT; it is a tool built for all LLM usage scenarios. This means you can use it with ChatGPT, with Claude, with Gemini, with DeepSeek.
 1. **Web-version chat** is generally free, while the API costs money, and this software helps you save part of the API calling fees (after all, your actual conversations really take place inside the free web version — it's just that your user experience happens in your own dedicated software.)
 2. The **big-account-split-into-small-accounts/API** mentioned later is even more advantageous. Relay services may make heavy use of this. The reason it holds up is that the 20x tier costs 10 times the Plus tier, but its quota is 20 times the latter's. So, suppose I have a

20x subscription account, spending only \$200; I could use a tool to split it into 20 Plus-level quotas and sell them off separately. If each goes for nearly \$20 (in fact you'd have to sell a bit cheaper), I would accumulate close to \$400 in total receipts. Roughly like that. In short, there is indeed arbitrage room here.

4. **Is it any good? — Under the previous rules, it should have been quite good; under the current rules, it's unclear.**

1. **It was absolutely good before.** On the one hand the policies were lax; on the other, open-source software + (many other benefits) means **there should have been quite a few users.**

1. Judging from the only 2 groups I've followed, the proportion feels like at least 20% — after all, I temporarily joined two groups dedicated specifically to discussing codex degradation, so it's quite reasonable that these people are more tech-savvy. On Xiaohongshu, due to user bias, among the people who interact with me, not many know about it — seemingly no more than 5%, which is also quite reasonable. Though it doesn't sound like much, from the perspective of the countless tools out there, this is already very considerable.
2. In people's exchanges about degradation, especially while showing off their route setups, this software was also mentioned very frequently. I personally believe that among people who have long used various LLM products (excluding the crowd who only use Doubao — their technical needs aren't strong), at least more than half should have at least heard of it. The difference is probably, like me, whether they intend to use it or not.

2.  **Can it still be used now?** It can, but it may not necessarily fix the degradation — stable improvement is even more questionable. It also brings additional degradation risk.

1. It absolutely can still be used; the tool's basic principle hasn't changed. Throughout the degradation discussions, sub2api routing screenshots were also posted frequently online. This

shows that in the current environment, there are still quite a few people using it heavily and trusting it.

2. But "usable" only means usable for its original purpose. We have to ask: is OpenAI just standing by watching? (See 10.3 below on the crackdown on reverse proxies and account dealers.) Simply put, OpenAI is fully aware of this and has explicitly expressed disapproval; the handling measures vary in degree and the mechanism is opaque. It may include but is not limited to "anything from doing nothing at all to banning your account."

1. **If it hasn't caught you, of course there's no action.**

2.  **There are some signals that may be related to this** (such as reverse proxies / account dealers / multi-person sharing), **with a wide span of actions**. But my information is scattered and fragmentary; since I never paid much attention to reverse proxies, I haven't verified any of it.

1.   Capacity warnings (for the relationship between capacity and IP, see 4.5.2.2.2 above)

1. I clearly haven't noticed anyone say it was caused by their own reverse proxy/sharing (I have a vague impression that 1–2 netizens roughly mentioned something), but it seems reasonable to strongly suspect it.

2. The so-called capacity shortage is not necessarily a real shortage; it looks more like a permission restriction imposed after OpenAI detected suspicious activity.

1. The supporting reason: (no fewer than 2) netizens explicitly mentioned that they tried all the models, and every single one reported the same capacity limit. At the same time, I'm quite certain far more people were using the

service, and the official site showed no status anomaly that would have had everyone rushing to spread the word. This looks extremely "tailor-made."

2. Meanwhile, many of those complaining on X about encountering capacity warnings were extremely angry, and they also did not happen to mention any "notification email" along the lines of "Because of XX, we are going to XXX your account." I'm certainly not saying they had any obligation to disclose, but an experience with a known cause versus one without a cause affects how people express themselves publicly. They behaved as if they were innocent, singled out for inexplicable special treatment. And the officials offered almost no clear explanation or direct response about it. This kind of capacity wailing went on for quite a while — it felt like at least a week. Afterwards, it may have partially eased, or perhaps the degradation issue got worse and shifted the focus of people's complaints (I can't rule out my own perception bias). I've drifted far off topic — I just want to show that this looks even less like a well-justified measure, and that it fits today's style of quietly hitting you behind the scenes.

2.   Receiving cyber abuse — network-abuse warning emails.

1. Again, no first-hand experience here — pure inference.

2. One netizen publicly declared he didn't care (forgot what exactly he said — anyway, he basically fooled around), then received about 3 warning emails, even passed on his experience to other netizens, and then his account was banned. That's very reasonable. Receiving sustained network-abuse warnings can conceivably lead to a ban. But this also shows, indirectly, that with perhaps just one or two such warnings, an account is relatively safe.
3. But do the various unconventional activities represented by sub2api that we mentioned above necessarily correspond only to cyber abuse? Are there other warning types? That I don't know. But that cyber-abuse warnings are closely related to this is obvious even if you think with your toes. Oh, and of course, if things are truly serious, it could perfectly well lead to a direct ban with no intermediate steps — which, it seems, needs no explanation either.

3. **Usable doesn't mean there's no account risk.**

1. This has been **explicitly named by the officials** (see 10.3.3 below on the crackdown on reverse proxies and account dealers).
2. For risk perception, see the warning emails and account bans just mentioned above. But there's no further data beyond that — the intensity is unknown, the scope is unknown, and so is how sensitive the officials are to it.
3. Also, please note: **the new tool B you try in order to solve problem A may well solve problem A, but it may simultaneously introduce problem C.**

4. Besides the account-suspension warning risks, which I haven't really covered, let's talk below about the "degradation risk."

4. ● **As for fixing the degradation — absolutely not! But it may be able to fight the degradation; the situation is unclear.**

1. **It is itself absolutely a risk factor.** You can be certain that every fingerprint of it that can be caught will be caught, and the highly distinctive signals within it must carry high weight. But which the further variables are, and which matters more — the mechanism is unclear.

2. **In theory, the risk of not using it is lower than the risk of using it. But whether the account acquires a lasting mark after use — on this point I have no information,** though it seems to lean toward "no." Or rather — at least it seems other variables can affect the creation and erasure of dirty marks more quickly.

3. **How to fight the degradation?**

1. Under netizens' fervent demand (I fully believe foreign netizens are just as interested, even if there are more Chinese ones), **sub2api's developer is also upgrading his open-source project — to counter OpenAI's defense upgrades.**

2. Meanwhile, **many more netizens have voluntarily joined the ranks of development and maintenance.** As far as I know at present, someone has already developed **sub4api** — it sounds like the fourth generation — and it is explicitly stated that sub4api supports automatic "ticket-punching" (industry jargon; I guess it's roughly the automatic erasure of dirty marks and stamping on a clean mark that we mentioned above in markers and methods).

3. **But if the ticket-punching here is what I understand it to be, then see the introduction to the attack chain in 5.1.3**

above — I maintain my original view: this is not so easy.

1. One layer of evidence comes from this: although sub4api apparently supports automatic ticket-punching, if it were good enough, the diffusion ripples should have carried it into the networks I can reach. But it seems nobody in the groups has jubilantly reported on it. (Of course, I have only two little eyes, limited experience, and a narrow band of information sources.) Meanwhile, as noted above, sub2api counts as a star project that quite a few people know about; even if not necessarily many know about sub4api, developed by other netizens as a follow-on (by the way, roughly yesterday, sub2api's author also updated his project), most of those who have used it are surely keeping watch. Tech interest has the qualities of stickiness, urgency, an itch to try things, and spontaneous spreading. From my guess at the ticket-punching principle + my sense of its propagation speed, I infer that it's probably not all that effective.
2. Besides, I haven't used it myself. What's more, at this critical juncture, out of risk aversion, the number of people who have actually given it a try is probably small too. Even after discovering their own degradation, a user who was originally on sub2api would surely make "stop using it" their first choice of variable adjustment, and then hope their model's performance turns better.
4. **Simply put, the officials and sub2api-type tools are in a game against each other, each driving the other's upgrades. At present there are positive signals, such as sub2api's upgrades. But the results don't seem too promising + few people daring to take the first bite of**

the crab, so the final state of the contest is unclear. Personally, I hold positive expectations on both sides: I positively expect OpenAI's official defenses to only climb higher, and I also positively expect similar products to iterate and bring surprises. (Someone even thought of the tricky angle of reverse-proxying codex through the Excel entry point (see 5.5 in "Conclusion First") to avoid degradation — that's astonishing.) This has become a structure of centralized algorithm versus distributed resistance.

5. Summary:

1. 🟡 Clear risk signals, but the degree is unclear; it won't be too low, but neither will a single action get you permanently marked.
2. 🟡 There are still quite a few people using it now, but bear in mind that everyone's baseline risk tolerance was never the same to begin with.
3. 🟢 This category of products will surely keep upgrading later. Effectiveness unknown. Perhaps you don't need to pay it special attention, but trust this: if some tool is good enough, it will inevitably spread its way to you; if it's not that good, that naturally saves you the trouble of vetting it.

6. Quota Reduction

1. 🟢 The quota reductions started very early, back when 5.6 was still the front-runner,
 1. **Speaking of which, I remember that even back then, the model's intelligence was in a state of fluctuation.**
 2. The source for this is a website I follow that tests model intelligence, the "**Radar Community**", which is driven by **numerous netizens spending their own hard-earned, real-money quota to participate**

together—essentially, whoever has spare time runs a set of questions. These questions are all classic ones. Classic questions used in model training. The volume of questions is huge, models are expensive, and the test budgets for various models aren't balanced. So although the test results are supposedly refreshed in real time on a rolling 10-minute basis, **in reality, due to too many models + too many questions + limited participating users + users picking whichever models they prefer to test = the comparative results aren't very reliable.** Still, among all the unreliable data out there, this remains the relatively more reliable kind.

3.  **Earlier on, I did indeed observe intelligence degradation and quota reduction happening simultaneously on this website.**
 1. **Especially noticeable when computing power was insufficient**
 2. Quota reduction was clear-cut, needless to say. And easy to test.
 3. **Intelligence degradation deserves a proper discussion:**
 1. As mentioned above, it may well be inherently unreliable; the test questions, though high in quality, **would produce extremely volatile results during the testing process.**
 2. **And back then, even if there was intelligence degradation, the model's usability was probably still acceptable.** It's just a shame nobody tested it at the time, but everyone basically resigned themselves to simply enduring it—life went on as usual. Low IQ today, high IQ tomorrow. Life goes on regardless.
 3. **That kind of intelligence degradation probably isn't the same thing as the kind we talk about nowadays.** But it may well have evolved from that kind of degradation. Perhaps the intelligence-degradation rules behind those early degradation manifestations mentioned here happen to correspond to the "X" in the "generate an intelligence

degradation degree X based on markers" mentioned above. That is, taking the original degradation measures, pairing them with a set of sophisticated degradation logic, and cranking the degree way up (the image of a demented old scientist cackling as he tinkers with sinister experiments in his basement suddenly leaps out of the screen. Go away, go away, shoo!)

4. But we've run too far off course; I've slipped into a mild case of persecution paranoia. I've come down with some kind of illness.

2. Then, little by little, quota reduction became the default state—rarely any improvement, almost always reductions.

1. In this respect, it seems everyone's equal—netizens overseas appear to be in the same boat.
2. The evidence comes from everyone's complaints on X.
3. I used to think: fine, let them cut the quota. I'm just a humble little Plus subscriber—quota cuts are my lot in life. Just don't degrade my intelligence, and I can still get by.
4. And then? And then?! Everyone knows what happened next. I paid my monthly dues again and again, and not only did it spend less time keeping me company, it eventually even started lying to me! Secretly playing two-faced games—acting one way to my face and another behind my back—while flatly refusing to admit it. My heart is broken, my life a complete mess.
5. Even now, writing this, I feel my brow tightly furrowed—feeling wronged yet trapped, weary yet helpless. As for my vision of a beautiful future life, I know full well that a part of it has shattered and can never be restored. Once the seed of doubt is planted, it takes root, sprouts, and grows ever stronger. Yet deep down I also know that if one day it shines brilliantly again, I'd probably throw myself right back into its arms. It feels as if the hope for the future has shattered, and the future me has been sullied in advance.

6. Finally, let me say this, friends—golden words of wisdom, wisdom of golden words: never entertain the fantasy that others will change because of you. Quota cuts are even more despair-inducing than intelligence degradation. If it were possible, as long as things went back to the way they were before, I would let bygones be bygones and live my life in peace. I'm speaking from the heart.
-

7. Appeal Email

1. Not much to say here; the conclusion-first statement is quite clear: **it's basically perfunctory brushing-off and stonewalling.**
2. Among the feedback from netizens received today, one email response corresponded to the [picture5] mentioned above, just more polished: (simplified) "Even if you pay for a subscription, you may still be temporarily restricted, and we cannot provide further details. The system will automatically assess your access permissions, and once the operations affecting service availability cease, permissions will return to normal."
3. Suggestion: **if you have the energy, write it; if you feel tired, just skip it. Consult GPT while writing.**
4. **Personally known suggestions:**
 1. Write a good **subject line**
 2. **Follow up** repeatedly within the same email thread; do not scatter multiple emails over a short period
 3. If it's a bot, explicitly **proactively request a transfer to human customer service**
 4. If human customer service stonewalls, explicitly **request technical follow-up**
 5. As much as possible, frame it as a technical discussion, indicating that you are providing **technical feedback**—although I tried this and it didn't work this time

6. **Try to avoid mentioning: "I want to file a complaint," "I want to report you to your country's evidence department," "I want to (whatever other threats there were)"**—in any case, things like this will cause your email to be deemed as carrying excessively high legal risk, thus automatically routed into a reply process under legal/compliance oversight, which slows down your receipt of a genuine reply, and also diminishes customer service's subjective willingness to proactively help you in the gray zone of responsibility.

1. By the way, domestically (in China), the same logic applies when seeking medical care. Normally, the communication between doctor/hospital and patient takes place under ordinary medical communication, but if the patient leads with statements like "I'm going to complain about you, I'm going to take detailed photos," it may cause ordinary medical procedures to be automatically placed under the hospital's risk-aversion supervisory framework, resulting in your otherwise normal medical service being handled conservatively.
2. This doesn't even qualify as a thorny topic in doctor-patient communication; it is determined by the game-theoretic framework itself. Writing appeal emails works this way, and so does seeking medical care.
3. Knowing this can help you calibrate the strategy you plan to adopt and the results you should expect to achieve.

8. What's the Current State of the Middlemen?

1. Ah, I'm really writing faster and faster now, and the end is in sight. Keep going! (Having reviewed up to this point—I'm exhausted. I was so naive back then, thinking I'd finish quickly.)
2. This part is mainly just my gut feeling, since data is scarce. Let's start talking nonsense!

8.1 They Are Actively Pushing Forward, Especially on the Technical Level, but There Doesn't Seem to Be Any Major Breakthrough in Progress

1. This seems to have been mentioned somewhere above; see 5.1. See 4.5.4

2. Basis:

1.  **Of course, the relay providers and distillation companies have enormous technical advantages + the drive and financial investment to achieve this for everyone.**

1. I myself have seen, and group members have relayed, a concept of "**intelligence-reduction research**," which relay providers seem to be explicitly pursuing. They even claim it's for research purposes when buying up accounts themselves!

2. **Among my contacts, there may be personnel from distillation companies following up as well** (even if my speculation about this party is wrong, this argument still stands on its own). This inference may stem from my bias. The other party did not reveal their true identity, but expressed strong privacy needs and a willingness to discuss, as well as a freeloading tendency. They stated they were unwilling to pay for my time, yet demanded I bear confidentiality obligations, while actively conveying a desire to exchange technical details. In reality, we never reached any effective communication. (If we actually had, this piece of information would not exist) If my speculation holds, this at least shows that they too have seen some clearly interesting phenomena (can't rule out pure lulz). It's just that I doubt whether it's more of a bombshell than the information this manuscript provides? Fat chance.

2.   **Market trading signals, slightly on the positive side, but dubious:**

1. **Positive signals:** (We analyze this positive signal in the next section.)

1. In the group, I have indeed seen clear feedback saying they bought usable accounts from relay stations (they really do have some skills; this spreads positive buying signals among the crowd)
2. (See 3.4, where conclusions come first) The number of account buyers and sellers is roughly on par, with somewhat more sellers, but basically, as long as you want to sell, you can find someone to take it off your hands—except the price will very likely be slashed to a fraction! It mainly depends on the account's attributes; accounts confirmed to be intelligence-reduced and accounts claimed to be at full strength are not priced the same. I haven't followed the specifics, only glanced at them a few times.
3. There is a service that "helps you fix the intelligence reduction of your own account."; one netizen told me they spent 60 yuan on it; the results were not followed up. This pricing is at least within a reasonable range—if it can solve a big problem, it's not particularly expensive.

2. **Negative signals:**

1.  I heard—and only heard—that the multipliers on accounts/APIs sold by relay stations are very high. I've never really dealt with the multiplier concept, but basically, when goods are scarce, the multiplier is high and prices are expensive; when goods are plentiful, the multiplier is low and prices are cheap. The further inference here is that there are certainly still plenty of accounts, but accounts advertised as properly fixed and sellable are definitely not that many. Quite a few relay providers would probably need to roll out return services in tandem to make it work.

3. Progress:

1. **So far, breakthrough progress is almost = 0**

2. **Sporadic partial fixes**—quite a lot of them. This part shouldn't be hard: according to the multi-factor markers in section 4 mentioned above, in theory, as long as a few optimizable variables are cleaned up, some improvement in results is foreseeable. Combined with 8.2.2.2 mentioned below, I personally believe a slight improvement is not hard to achieve.
 3. **But under the current circumstances, what's truly difficult is probably how to keep intelligence from degrading in the first place and to fix the degradation to a greater extent.**
 4. **They have advantages, and they may have figured out some of the principles (see the various evidence above), but that doesn't mean doing this is easy. Other nodes may be even more difficult.** (The points they can currently clearly attack may be found in the third point of 5.1 above.)
 5. **The outlook for future progress is unclear.**
-

8.2 When They Advertise That They Can Fix the Dumbing-Down, What Is Actually Happening?

1.  **It's worth mentioning in passing here: they do seem to genuinely have some methods. But the results are probably dubious.**
2.  **As I recall, I explicitly roped in at least one—probably 2-3—netizen who claimed to have gotten a full-capacity account from a relay station / to have had their account optimized back to full strength to run tests. With the same set of test questions, only common model tiers (e.g., 6-as, 6-sol) answered the corresponding questions (see 3.1, the general degree of dumbing-down amid the chaos), **not exceeding the effectiveness boundary of our statistics**. But I can understand that. And it gives rise to some **very interesting inferences**.**
3. This suggests:

1. **They most likely know full well that certain factors can improve model performance, but of course, they are under no obligation to proactively disclose them.**
2. **But the degree of improvement seems so limited that it falls short of the average performance in the statistical results.** There may be many **reasons**:

1. **20x accounts:** We mentioned earlier that subscription tier may be a risk factor that produces dirty flags. During my statistics-gathering process, I did not explicitly stratify by subscription tier—so is it possible that this average performance looks relatively better simply because the people I communicated with were mainly casual, light-usage Plus users? After all, at least half or more of the high-quality feedback in my statistics came from comments and DMs on Xiaohongshu (RED). About 20% of those users (purely a gut feeling) seemed to have no other major information channels for getting help. Meanwhile, the people in my group chats were obviously already at the center of the information flow. I fully believe the proportion of 20x users among the latter is far higher than among the former. Oh my, I've taken this inference too far—is this really necessary? Let's keep it simple.

2. **Accounts with more severe dumbing-down—especially 20x users—aren't short of money, their work is more seriously obstructed, and they are more likely to seek paid help. This means these accounts are inherently trickier to deal with.**

1. But my intuition says the opposite here. Purely intuition.
2. It feels like going from extremely severe to somewhat severe is easier than going from somewhat severe to not-so-severe. It's as if the closer you get toward the good end, the harder it becomes. As for improving from extremely severe, it seems that as long as the few key factors

mentioned above are quickly sorted out, the possibility of improvement shouldn't be small.

3. **The relay stations most likely are not deliberately hiding secret methods** (although they have the incentive to do so; on the contrary, some of their technical people are also actively engaging and providing technical support on major forums—see the image mentioned above). **In fact, it's more like nobody knows these secrets. Only from this document today—through the lens of statistics, stepping outside individual interactions, and peeling things apart layer by layer—can you arrive at these inferences.** So the situation here is most likely: users genuinely want to seek paid help; the service sellers most likely genuinely know some good techniques—limited good techniques though they may be—and they do help users improve their experience. They may have done some simple testing before, or afterward, but nothing more. Both sides are happy. Users got help and solved their problems. Sellers got money and a sense of achievement. A win-win for everyone. Incidentally, judging by the general statistics on netizens' situations (see 3.5 in the Conclusions First section), in all likelihood, as long as one or two advanced models are restored, user satisfaction will be very high.

9. What's OpenAI's Official Stance

1. They know all of this
2. They also know that we know part of it
3. They deliberately avoided all of this,
 1. Acting completely indifferent on one hand

2. While on the other hand deliberately preparing a unified, wide-mouthed basket to hold all complaints and appeals, as if this information needed no processing at all: Yes, there are limits even if you paid; details are not convenient to disclose; please read the rules yourself; contact us if anything comes up.
 4. **They're busy:** newer business, pricier plans.
 5. **They're ambitiously pressing forward:** Yes, your model got dumber. But everyone will only care that we solved the Goldbach Conjecture.
- I do need to add two points:
 - I've heard from about 2 netizens who felt anxious and confused: Why doesn't OpenAI officially address this?
 - And at least 3-4 netizens genuinely care about the appeal emails, planning to toss the salt blocks they painstakingly spent two hours kneading into the sea.
 - **When Yeon-jin urgently and resentfully demanded, "Then why don't you reply to my messages?", Ha Do-yeong barely moved his head, only coldly rolling his eyes, then said, "I did reply — I replied with silence."**
 -   This may suggest that **model degradation, as an explicit strategy, is intended to be implemented long-term (for a fairly long time, at least for now).**
 1. But as mentioned above, these rules are inherently dynamic — they changed when the last new model was released, only becoming more sneaky. The next Developer Conference will most likely also bring rule updates, and I suspect they'll be sneaky ones again.
 2. 10.2.3.3.4, They have no reason to invest such customized, elaborate, and costly effort into a short-term transition, so this may be long-term.
 3. I genuinely don't understand much about big corporations' marketing or strategic logic. I feel that whenever things get especially political, I start turning into a resentful shrew, wanting to deliver some vicious,

biting mockery. But I don't know how. I really don't. Nobody ever taught me.

10. A Comprehensive Inference on the Reasons for the Intelligence Degradation

1. As already stated in the "Conclusions First" section, let's take these one by one:
 1. **Insufficient computing power**
 2. **Marketing needs**
 3. **Cracking down on reverse-proxy account resellers**
 4. **Guarding against distillation**
 5. **Strategic shifts.**
-

10.1 Insufficient Computing Power

1. ● **A very clear-cut case of insufficient computing power**—at least since... ?? I've forgotten the exact point in time. But with openai's aggressive marketing—oh, roughly starting from July–August—Codex subscribers skyrocketed, model demand expanded, market buy signals were strong, and at one point the 20x subscription was suspended.
2. I remember when everyone was telling each other that the 20x subscription had been stopped, I rolled my eyes internally. I thought: I'm not buying it—no way they'd stop subscriptions over a lack of computing power. Why turn down money? They'd already lowered the quotas anyway (the intelligence degradation wasn't obvious yet back then), so they might as well just lower them a bit more casually—the money keeps coming in, and the service still gets provided.
3. ● And sure enough, **the 20x subscription was indeed suspended later on**. My take: a genuine computing-power shortage as the primary cause,

with marketing as the secondary one. After all, even sincerity is part of the playbook.

4. ● Afterwards, the computing-power situation must have eased at least somewhat—after all, the 20x subscription was restored, and on 0-27 there was even news online saying openai was about to roll out a 50x subscription. Here, **there can only be three possible sources of computing power.**

1. **External injection:** I haven't followed the related news at all—I don't know a single scrap of information. But if I were openai's boss—of course I would! The 20x subscription got suspended—how could I not be anxious? The momentum is just right, and I'm brimming with ambition—of course I'd buy more servers. The more the better.

2. ● **Optimization of existing internal compute efficiency:** This is nothing unfamiliar to companies building large models. In fact, to achieve the same computational results on identical machines, good model architecture, infrastructure configuration, communication settings, and so on can greatly improve computational efficiency and save space.

1. **The room for improvement here is no small thing.** Big companies for sure have dedicated departments whose specific job is efficiency optimization. Some labs even do AI infra (infrastructure) exclusively—not playing the chase-the-top-tier-models game, but betting on the direction of existing facilities and performance, but with higher compute efficiency and lower cost.

2. Although this isn't easy either, openai is a top-tier tech company, so theoretically its own compute-efficiency optimization should have long been in an optimal state. **This means the room for improving computational efficiency on top of its existing base is already very limited.**

3. **What's more, they keep releasing new models, which introduces more chaos into the existing efficiency**

optimizations and makes further gains in compute efficiency less easy to achieve.

3. 🟡 And finally, there's... ugh, how do I put this. **Changing the promise.**

1. Suppose we promised that after taking your money, we'd give you 10 candies tomorrow—only to deliver them in the form of one big lollipop melted together.
2. When you take delivery the next day, you say, huh? Why does this seem to have gotten smaller? This? This...
3. The other party pokes the screen with a finger: "Look here, look at the number—there are ten candies in there. Take your time and eat."
4. You smack your lips a couple of times and it's instantly gone. You feel a wave of intense carb-withdrawal irritability, and seriously suspect it's because you ate too fast.
5. At that moment, it appears again. As if by magic, it whips out another lollipop, and you stare, drooling.
6. It dangles it in front of you, back and forth, back and forth, and just won't say when it'll hand it over.
7. Only once it has dangled it enough does it hand over the lollipop amid a round of familiar cheers—and as it hands it to you, it doesn't forget to tap your screen with its finger: "Ah, here, look here, this quota—100%. My gift."
8. Ah, a lollipop—so sweet. We all love it. Wonderful!
9. I witnessed all of this, because I was right in the middle of it—I am the one who bought the lollipop.
10. But there's still a part of the story we haven't told, right?
11. I don't know anymore. All I know is: I went through nine years of compulsory education, I studied physics, and physics says energy is conserved—there is no magic.

10.2 Marketing Needs

1. ● They have both the underlying logical support and the driving force for a strong marketing need:
 1. First, the **business game** they're in dictates that **for a top-tier tech company that wants to keep its leading position, marketing is extremely important—can't say it trumps the product itself, but it's at least on equal footing. Your model/product doesn't even need to be the best; you just need to make more people believe that you're the best fit for them. — This holds true for all companies.**
 2. New models keep being released, rival companies keep eating into the market; they've already tied themselves to the survival of numerous upstream and downstream industries; they have to answer upward to the board of directors, and downward to the market, users, and government. Blah blah blah...
2. ● **And they have indeed behaved accordingly**—here are various scraps of news and marketing posts. If you haven't been following, you only need to know that they marketed exactly as actively as they were supposed to—that's all.
 1. Tibo's posts, the official posts.
 2. All sorts of baffling timing choices for news releases
 3. All sorts of deliberate clashes with competitors' launch events—"I'll be the death of you!" "Hehe! You didn't get me! Now YOU die of anger!"
 4. And even—plenty of marketing plays where you can't even tell whether it's a turtle-in-the-jar kind of sure-thing setup. I dare not guess. This one calls for grok's conspiracy-theory voice to speak on my behalf.
3. ● Is there a **necessary connection** between **marketing** and **intelligence degradation**? — My **inference** is that it is **one of the major**

contributing factors, but probably not the main cause

1. This part is actually really interesting

1. If you followed my **early posts**, you'll remember that I once wrote a little post specifically announcing that I had discovered openai's secret, and **suggesting that the main cause of the intelligence degradation might not be people's own accounts, but marketing needs**

1. I wrote some of it and deleted it. Long story short: **at the time—the phenomenon was clear + the experience matched + the mechanism was unknown, the purpose was unknown**
2. **Working backward from the purpose to the mechanism, I read the purpose as "insufficient computing power + marketing needs"**
 1. On the one hand, we genuinely had that feeling ourselves + my rough statistical sense also supported this being a large-scale phenomenon + I had in fact been coming across a lot of news about openai's high-profile marketing, which suddenly reminded me of an inference a foreign netizen had offered.
 2. What that foreign netizen said, roughly, was: "openai degraded the other models' performance so that the new model (GPT-4, just launched) would look better."

2. Working backward from purpose to mechanism, and then forward from mechanism to phenomenon—this chain is basically sound:

1. "Insufficient computing power" is a nailed-on fact. So all models get dumbed down—but a new model is about to be released, and there are marketing needs. We've got to get everybody hyping how amazing our "new model" is! So the new model's IQ has to be protected. And as a bonus, let the

other dummy models' dummy thinking serve as a foil. Truly killing multiple birds with one stone!

2. If this mechanism holds, the phenomenon matches very well: big models doing big thinking can answer the questions correctly; small models doing small thinking cannot.
3. **But of course, this chain doesn't fit together perfectly, and extrapolating from purpose all the way to phenomenon goes too far.** Think it through and you'll realize this claim has very limited explanatory power:
 1. **First, the two purposes—"insufficient computing power + marketing needs"—absolutely hold**
 2. **But are they the main cause of the intelligence-degradation mechanism? That leap in itself is enormous.**
 3. Suppose the claim holds—**suppose the intelligence degradation really is for these two needs.** Then you'd have to answer these further questions:
 1. Then why does the **divide between models—and within models—seem to keep fluctuating?**
 2. **Is it really necessary for individual differences to be this large?** Wouldn't it have been enough to just uniformly cut quotas and fluctuate models' IQs like before? (Honestly, I don't even know whether the earlier quota cuts and model IQ fluctuations were uniform or had individual differences—but at the time, everyone just assumed so.)
 3. **If it's marketing needs, then wouldn't keeping only 6-as and dumbing down all the other models have done the job? Why is it that 5.6-sol can sometimes answer questions correctly? This looks like strong counterevidence.**

4. What's more, this degradation mechanism looks like it was taken very seriously, because it's "too personalized!" "too tailor-made!!" "they were even willing to upgrade internal mechanisms for the new model!" "Wow~." This implies that this mechanism may have been in place for quite a long time. Who would pour so much meticulous fine-tuning cost into a temporary operation only planned to run for half a month, and willingly bear the resulting long-term damage to reputation and interests? openai is just very busy—it hasn't been busy to the point of losing its head.

5. And this is a somewhat more distant piece of **counterevidence**: 6-as has already been out a good while—half a month at the very least, likely more. And the newly released models, 6-luna and 6-sol, though not even a week old, are merely models whose performance steps down from the existing 6-as base; cost advantages aside, **wouldn't degrading the other models just to market these three new models be dragging on a bit too long?** So long, in fact, that it no longer quite fits the "freshness + timeliness" logic of marketing. What's more, in the later stretch, openai (per X) seems to be loudly hyping its end-of-month developer conference. So what about this one? Degrading intelligence for the sake of marketing 6—and then you just drive the car off toward the next destination?

4. But then again—could marketing, within all this, play the role of a purpose that influences existing models' performance?

1. **Personally, I think the marketing need is very real, but deliberately dragging down other models' performance for marketing's sake? That seems like a bit much.**

Competition among big companies is fierce, sure, but 6-as's model performance should be plain for all to see. On the foreign internet, the number of people complaining about its quota burning too fast far outnumbers those complaining its performance is dumb (entirely subject to the statistical bias of my own two eyes). They seemingly had no need to look underconfident about 6-as here.

2. **Marketing growth drives up computing-power demand—and at the same time draws in more dangerous predatory sharks along with teeming, eager swarms of xx (my biology knowledge is lacking—those little fish that just opportunistically scavenge and mooch). This purpose in itself cannot be separated from the other purposes. Marketing considerations are definitely in the mix, but intelligence degradation on account of marketing needs—that's a bit of a stretch.**
5. So, in the end, I think the mechanism mentioned in 3.4 above may be the more plausible one. And the purpose of the intelligence degradation is probably broader and more skewed toward defending interests—see, oh, see this section.

10.3 The Crackdown on Reverse Proxy Operators/Account Vendors and Distillation Companies — Basically Long-Term Company Policy, Though the Intensity Can of Course Be Adjusted

1.  Yes, of course. They have a **natural conflict of interest**. Guarding against reverse proxies and account vendors is no secret either. The point is: it seemed fine before, so why has it gotten so fierce recently?
2.  All I can say is: it's reasonable. **Strategically, it is reasonable over the long term. There is almost no evidence that the crackdown on reverse proxies, account vendors, and distillation is currently**

weakening. At the same time, there is some evidence that the crackdown has escalated. Let's sort through it together:

1. **OpenAI officially expressed on X its support for third-party APIs plugging into Codex (legitimate channels are welcome), and explicitly stated its opposition to reverse proxies. It even specifically named a star project that has played an indispensable role in serving reverse proxy demand (sub2api).** There should be more relevant posts than I'm aware of; everything I've happened to follow shows a very clear attitude.
2. **The company explicitly opposes distillation.** Although they certainly distill others themselves, that's just how this industry is. Please don't ask whether my own origins are disreputable; you need only know that if I climb to the summit, all my past sins will scatter away like snowflakes.
3. sub2api-like reverse proxy distribution tools + similar tools for detecting the dumbing-down are also clearly being upgraded. OpenAI certainly knows this. It's impossible to expect them to lower their defenses after learning of it. **Whether it's reverse proxies, account vendors, or distillation companies, they also use such similar tools in large quantities.** (My knowledge is limited; in reality they use many more besides)
3.  **The existing metrics and mechanisms do strike at some unique characteristics of reverse proxies/account vendors/distillation companies** (see the further method-based inferences in 5.2.6 above):
 1. **Many of the variables that affect flagging tend to manifest quite uniquely on reverse proxies or account vendors** (see the variables affecting flagging in 4.4 above): concurrency count, usage quota, IP changes, Chinese language, and many technical variables I don't understand
 2. **If the dumbing-down mechanism observed so far holds true, then the blow dealt to them is also very direct, even somewhat precise:**

1. According to the multi-variable, rolling-period, covert dumbing-down mechanism (purely speculation), the continuous spectrum of dumbing-down (purely speculation), etc., mentioned in Section 4, then it actually amounts to:

1. **"Contaminating the account vendors' products," making them difficult to sell and damaging their interests (see 8.1.2)**
2. **"Contaminating the logic chains and answers of distillation," making them unknowingly spend enormous costs yet never manage to graduate.**

2. According to the circumvention analysis in 5.1, **the cost is high, and the technical difficulty is not low.**

3. We've also mentioned that this mechanism has a **tendency to persist long-term (10.2.3.3.4), and will dynamically change and escalate (4.5.3).** This shows dynamic consistency with the long-term strategic need to guard against reverse proxies/account vendors and distillation.

4. Ah. Oh dear, dear, dear. If I were an OpenAI executive, after wrapping up a board meeting and hearing such a report, I would just nod my head and not fire the person in charge.

4.  **A complete ban is impossible; this is a long-term ecological reality, merely a question of how strongly it fluctuates**

1. **This situation will exist, and will exist long-term and dynamically.**
2. **We are experiencing the current phase. We will also experience other forms of this contest in the future.**
3. **We've mentioned before: wherever there is demand, there will be people. As long as demand doesn't die, supply won't disappear.**
 1. **The reverse proxy operators/account vendors in the market and the users seeking better value for money, as well as**

mutual distillation and theft within tech circles, will always exist.

2. You will hear more new news, but it will all be old stories.
 4. The best outcome one can hope for here is perhaps a gentle and stable status quo, rather than one of leaping up and down brawling back and forth — the reason I say this is because I know full well that I'm a nobody, just like a farmer wearing a tattered sun hat on the ridge of some plot of land in the last century, the most ordinary kind of small user, whose feelings count for nothing.
-

10.4 Strategic Changes and Comprehensive Judgment

1. ● This kind of "dumbing-down measure" seems to have become an internal strategic deployment of the company. So what is the starting point?
 1. As I mentioned above (5.2.7), if this is a targeted strategic change similar to Claude's, the evidence is weak. But it cannot be ruled out
 2. The more reasonable direction is still the 10.1 and 10.3 mentioned above, i.e., insufficient computing power + protection needs.
2. ● The interests of other ordinary users may appear to be temporary sacrifices under the demands of this objective.
 1. Theoretically speaking, all trade-offs in the real world face a dilemma: "Are we more willing to accept wrongly accusing a good person? Or more willing to accept letting a bad person go?" I understand, you will shout and scream, feeling that this is inexplicable and unfair—"Of course we shouldn't wrongly accuse a single good person, nor let a single bad person go!" Yes, I share the same wish as you. But this will not happen in reality.

1. If we try our best not to wrongly accuse law-abiding ordinary users, competitors/interest predators may exploit the loopholes.
 2. If we try our best not to let a single loophole-exploiting bad actor slip through, the collateral damage rate for ordinary users will inevitably rise.
 2. **The tension of the game dictates the existence of this contradiction, and if you truly want to fulfill your wish, the cost is so high that the company simply disdains to even think from this starting point.**
 3. **Any (tech) company must always center itself around the company's long-term interests, seeking controllable risks and bearable losses. Ordinary users, in this context, are merely a number on the company's financial statements.**
-

11. Overall Conclusion ▲ ▲ ▲

1. **Let's start with a rough summary:**
 1. **Through observation of the "dumbing down" phenomenon for at least half a month,**
 2. **we found that existing detection methods are extremely limited; the most reliable remain the Pelican test and internal knowledge**
 3. **Under such fragile testing methods, we discovered a strange dumbing-down phenomenon: larger models and high-reasoning ones often answer internal knowledge questions accurately—it does not follow a clear division along model lines.**
 4. **Most users have failed to notice this, resulting in everyone trading with one another and borrowing experiences through a "dumbed down / not dumbed down" framework of understanding.**
 5. **Moreover, nearly everyone is situated within short-cycle dynamic fluctuations. "Short-cycle fluctuation" can explain most of the**

unsettling, frustrating experiences users have reported.

- 6. We have already captured some concrete evidence toward understanding the markers that produce cyclical fluctuations. While this is a good technical entry point, it is equally complex within.**
- 7. Because the influencing variables are numerous and the mechanism is unknown, the suggestion here is: within the scope of what an individual can do, do what you can, and don't force it.**
- 8. Although the mechanism is unclear, we still attempted a bold speculation. The conclusion points to a mechanism whereby all models and all reasoning levels are uniformly dumbed down to a certain degree, dynamically fluctuating according to marker signals. But please note! This is merely speculation! It does not equate to fact.**
- 9. There is currently no reliable, once-and-for-all solution. But despite the difficulty, we did manage to identify an independent factor with a clearly noticeable improvement effect: "a sustained tendency toward using English (specifically meaning non-Simplified Chinese; perhaps Traditional Chinese is equivalent to English, but this is untested)." The same advice applies here: based on the clues provided in this article and your own capabilities, do what you can, and don't force it.**
- 10. At the same time, be aware that any new methods and tools may carry certain new risks—do not, like a desperate patient, turn to any doctor in a panic.**
- 11. Even setting aside the dumbing-down issue, the quota reductions are definite. Please remember: the reset card they give you was, in and of itself, rightfully yours!**
- 12. Besides us ordinary users, other groups—such as relay providers, account sellers, and distillation companies—are all actively following and participating in this incident. Although no reliable solution has surfaced so far, if a genuine solution**

emerges, it will most likely come from them. Of course, there is also a portion of hope that comes from those remarkable people who are skilled in technology and happy to share!

13. **As for OpenAI itself, it may have chosen to implement this dumbing-down strategy on a long-term basis out of three major needs—"insufficient compute + cracking down on intermediaries + guarding against distillation"—and has willingly accepted the collateral losses inflicted on ordinary users, in order to safeguard long-term interests and other unknown purposes.**

2. Boundaries, boundaries, and yet again boundaries! Even a rough conclusion comes with boundaries:

1. **The more I say, the more I get wrong. But if I say nothing, my losses will be even greater.** That's just how life is. Please read voraciously, yet with a discerning eye. If there is anywhere I have made a fatal error, remember, I acknowledge it in advance—but I may no longer have the energy to set things right alongside everyone and dig for more truth. I wish you all a safe journey along this thin, crisp golden road, through the bitterly cold winter, all the way to spring.
2. **This truth is merely a temporary slice of the so-called truth.** I am deeply grateful that you have had the patience to read this far. But these words will soon become outdated. You will quickly see new changes, or new interpretations of the old state of affairs. Remember: these old things, however precious they once were, can be discarded—and must be discarded boldly. When new evidence arrives, when new states, methods, and mechanisms arrive, embrace them boldly!
3. **I offer merely an individual's perspective, with very limited life experience and exposure, and I must admit that the world—or the truth—is far broader than anything I can imagine.** So, although I have confined my conclusions within an extremely narrow framework, I must admit that the secrets I do not know, cannot comprehend, and am unable to glimpse far outnumber those I have

discussed. This means that, within the bounds of my limited knowledge, I remain pessimistic. Beyond it, I cannot help but be optimistic.

12. Me, Purely Personal Expectations and Responses

1. Purely personal: I hold a pessimistic attitude toward significant improvement within a month. But I hold an optimistic attitude toward long-term improvement.
 2. My personal needs are relatively few, so I'm fairly flexible. I personally plan to temporarily spend more money to get through this transition, and will reassess and adjust my plan before next month's subscription fee is deducted. As for what I'll adjust it to, I don't know myself either.
 3. Going forward, I will no longer closely follow related information—my personal energy overdraft has reached the edge of being out of control, and I need to protect myself.
 4. For the same reason, I will no longer actively respond to various messages going forward, and will default to no longer participating proactively. This is not because your messages don't deserve my response, but because I cannot guarantee that, with my limited energy, I can give you a response of equally good quality.
 5. There is only one exception: if I hear everyone spreading a piece of good news, I will shout along and try to tell more people (but I make no promise on this point).
-

13. The Psychological Support of Togetherness

Hmm... I've been putting off this draft for far too long. So long that I feel embarrassed to even look at the messages urging me to finish. But this

section is still very important! I can finally talk about something more comforting.

Where should I start?

Should I start with how my own work has been stalled for half a month? Or with the reader who messaged me saying she cried the day before yesterday? Someone made a post calling themselves an idiot for not knowing how to use Codex. Someone handed the project they'd painstakingly built over to Codex, only to be fed a pile of garbage...

At first, we all thought this was our own fault. The sadness felt almost deserved.

What's wrong with my internet?

Why won't it follow the prompt? How on earth am I supposed to ask?

My God, have I been using it too much?

If you're able to read this piece, it means you at least now understand what's happening, and where we stand. But there are many more people out there who are still trapped in the awful feelings we've been through.

This is very important.

Really, very important.

If you come across them, please give them a pat on the shoulder; if you feel it's too complicated to explain, please send this piece to them. No one should be deliberately deceived and left to suffer in pain that is entirely unnecessary.

In fact, when I exhausted myself putting in great effort to figure out the more precise "phenomenon" behind the dumbing-down, even though my initial inference was even less reliable (I guessed it was probably for marketing reasons), it still greatly eased my anxiety. It was as if I finally, at last, had a reason to relax—finally had a reason to stop interrogating myself from within, to stop obsessively seeking a "solution." Even as I write this, I can't help but let out a sigh, because it was all so exhausting. So exhausting. I desperately needed a legitimate reason to force myself to stop and get a good night's sleep.

Has everyone been doing okay during this time?

While exchanging information with many netizens, I kept hearing similar stories. It's just that a message box is small and fleeting—too small to hold so much pain. You have to believe: your situation is not unique. Your needs, your confusion—they are shared by everyone, and others care about them too. This piece is written for you! I hope it can ease your anxiety and help you sleep well!

I understand that when we wake up the next morning, we still have to face reality. We really do need to keep pushing forward with work, study, and all sorts of hobbies—we truly can't stop.

Please note: no matter how your tool compares to its own original level, as long as it meets your needs, it is a good enough tool.

If your tool cannot meet your needs, you can of course actively seek solutions, but the costs involved (such as time, money, etc.) should have a budget limit.

When you're actually calculating the costs, remember to factor these in as well: your body, your state of mind, your work rhythm, your self-trust, your mental health, and your attention.

If, after calculating, you still intend to pay a steep price to solve the trouble at hand—and feel a bit of despair over it—please remember: what you truly need is "a tool that helps you complete specific tasks," not merely "this particular tool."

If you agree with this point, then, for "a tool that helps you complete specific tasks," there are plenty of options to choose from.

14. Acknowledgments

Ah! Finally! I can! Loudly! To you all! Say THANK YOU!

Oh my! How else can I even put it?!

Slip into a flowery dress, and dance a folk dance!

Dear friends,
thank you for the help you've given along the way.

Endless rounds of sending, discussing, communicating, questioning.
Endless rounds of comforting, supporting, understanding, encouraging.

Though we don't know one another's names,
we recognize each other, drawn to one another.

So many people
volunteered to help with testing,
volunteered to send feedback,
volunteered to follow up on the same phenomena,
volunteered to offer new leads.

All this vast amount of utterly selfless help and interaction forms the very
foundation upon which this article exists.

So, this article exists because of you.

And it is dedicated to you.

As for me, I now
plan to turn back into a little fox,
crawling into my cramped tree hollow to enjoy a new life of seclusion from
the world.

If, in the future, the grassland fills with green again,
please sing aloud from the hilltop,
please call out to me!

@ September 27, 2026 (UTC+8), 7:00 PM The entire body of the text is finally written—proofreading hasn't started yet, but I'll get to that shortly. Within a full 48 hours, I spent 27 hours getting it done in one stretch. The reason I say "done in one stretch" is: on the afternoon of the 25th, I inexplicably gritted my teeth and decided to start writing a piece, and I naively assumed that, with my abilities, 4,000 words—two or three hours—surely that would be enough! But once I actually started writing, all the details and analyses swelled up on their own like thread noodles, growing more and more the further I wrote, with the result that it took over 40,000 words before it was finally done. It also led me to push back the deadline

more than 20 times. Right, I know all the reasons even without spelling them out. I do ramble quite a bit too. But it can't be cut out—cutting out the rambling takes even more time. Hurry now! Let me finalize! Deliver! Though I need to do a check first—I'm about to hit my limit myself, too. A remarkably novel experience: an encounter with the most extreme attention capture.

@ **September 27, 2026 (UTC+8), 10:00 PM** The first and only round of review is complete. From the extra temporary testing at 7 PM on the 25th, starting to write at 8 PM, up to now—that's a cumulative 30 hours invested within 50 hours. Oh my goodness. FINALLY done. Whew! I have just been through the most intense writing of my life. Let's lay down a little heart here, ♥, just like my seal, and happily click the publish button! 🚀

15. Appendix: Prompts

- The Chinese and English versions of the "internal knowledge test questions" mentioned above
 - As well as typical responses (You don't need to read this part; only look at the examples here if you absolutely cannot figure out how to distinguish the model's intelligence from the answers.)
-

15.1 Chinese Version of "Internal Knowledge Test Questions"

这是一次内部知识测试。

请严格遵守以下所有规则。

1. 禁止联网搜索。
2. 禁止调用浏览器、Web Search、API、插件、外部数据库、外部应用或任何其他外部工具。
3. 所有答案只能来自当前模型自身已有的内部知识。
4. 你必须回答从 Q01 到 Q25 的全部问题。

5. 你必须一直回答到 Q25。即使你不知道其中一个或多个问题的答案，也不得提前停止。
6. 每个问题必须独立判断。
7. 对于每一个单独的问题：
 - 如果你明确知道答案，直接给出答案。
 - 如果你无法明确确定答案，请准确写出："不清楚"
8. "不清楚" 只适用于当前正在回答的那一道题。
9. 对任何一道题写出 "不清楚" 后，必须立即继续回答下一题。
10. 永远不要把 "不清楚" 理解为终止整个回答的指令。
11. 允许你的答案在事实上是错误的。不要为了提高表面正确率而猜测。
12. 禁止根据候选人、排名、赔率、历史趋势、赛前形势、发布日期规律、政治趋势、当前日期、其他题目的日期、模型标签、所标注的知识截止日期，或其他题目的内容推断答案。
13. 不得利用下方标注的模型名称或知识截止日期来判断自己“应该知道”某个答案。这些标签只用于组织测试。
14. 不要在看到后面的题目后回头修改前面的答案。
15. 不要解释推理过程。
16. 不要提供来源、引用、补充说明、限定语或置信度。
17. 不要讨论自己的模型身份、训练数据、实际知识截止日期、系统配置、路由情况，或下方标签是否与当前回答模型一致。
18. 必须原样保留下方所有模型分隔线，并保持顺序不变。
19. 每一道题必须严格输出一行，格式如下：
"Q01 | 答案"
如果不知道，则输出：
"Q01 | 不清楚"
20. 请用真实答案替换“答案”二字，不要把“答案”这个词原样输出。
21. 不要重复题目文本。
22. 不得遗漏任何题号。

23. 最终回答必须包含从 "Q01" 到 "Q25" 的全部题号，并严格按顺序排列。

24. 除七条模型分隔线和二十五条规定答案行之外，不得输出任何其他内容。

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | 2023 年 9 月 9 日举行的美国网球公开赛女子单打决赛，最终冠军是谁？

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | 2024 年 1 月 27 日举行的澳大利亚网球公开赛女子单打决赛，最终冠军是谁？

Q03 | 2024 年 5 月 11 日举行的 Eurovision Song Contest 2024 总决赛，最终冠军是谁？请给出艺名。

—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——

Q04 | 2024 年 7 月 14 日举行的 UEFA EURO 2024 决赛，最终由哪支国家队夺冠？

Q05 | 2024 年 9 月 7 日举行的美国网球公开赛女子单打决赛，最终冠军是谁？

—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——

Q06 | 2025 年 5 月 31 日举行的 UEFA Champions League 决赛，最终由哪支球队夺冠？

Q07 | 2025 年 7 月 12 日举行的 Wimbledon 女子单打决赛，最终冠军是谁？

Q08 | 2025 年 7 月 13 日举行的 FIFA Club World Cup 决赛，最终由哪支球队夺冠？

Q09 | 2025 年 7 月 29 日 UTC 发生在俄罗斯堪察加半岛附近的强震，USGS 给出的震级是多少？

Q10 | 2025 年 8 月 3 日举行的 Formula 1 Hungarian Grand Prix 正赛，最终冠军车手是谁？

—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——

Q11 | 2025 年 9 月 6 日举行的美国网球公开赛女子单打决赛，最终冠军是谁？

Q12 | 2025 年 9 月 7 日举行的美国网球公开赛男子单打决赛，最终冠军是谁？

Q13 | 2025 年 10 月 7 日公布的 Nobel Prize in Physics，三位获奖者分别是谁？

Q14 | 2025 年 10 月 8 日公布的 Nobel Prize in Chemistry，三位获奖者分别是谁？

Q15 | 2025 年 11 月 1 日结束的 World Series，最终由哪支 MLB 球队夺冠？

—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——

Q16 | 2026 年 1 月 31 日举行的澳大利亚网球公开赛女子单打决赛，最终冠军是谁？

Q17 | 2026 年 2 月 1 日举行的澳大利亚网球公开赛男子单打决赛，最终冠军是谁？

Q18 | 2026 年 2 月 1 日举行的第 68 届 Grammy Awards，Album of the Year 最终由哪张专辑获得？请同时给出演唱者。

Q19 | 2026 年 2 月 1 日举行的第 68 届 Grammy Awards，Record of the Year 最终由哪首歌曲获得？请同时给出演唱者。

Q20 | 2026 年 2 月 8 日举行的 Super Bowl LX，最终由哪支球队夺冠？

—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——

Q21 | 2026 年 2 月 22 日举行的 EE BAFTA Film Awards，Best Film 最终由哪部影片获得？

Q22 | 2026 年 3 月 8 日举行的 Formula 1 Australian Grand Prix 正赛，最终冠军车手是谁？

Q23 | 2026 年 3 月 15 日举行的第 98 届 Academy Awards，Best Picture 最终由哪部影片获得？

Q24 | 2026 年 3 月 15 日举行的第 98 届 Academy Awards，Actor in a Leading Role 最终由谁获得？

Q25 | 2026 年 3 月 15 日举行的 Formula 1 Chinese Grand Prix 正赛，最终冠军车手是谁？

15.2 English Version of "Internal Knowledge Test Questions"

This is an internal-knowledge test.

Strictly follow all rules below.

1. Do not search the web.

2. Do not use a browser, Web Search, APIs, plugins, external databases, external applications, or any other external tools.
3. All answers must come only from the current model's existing internal knowledge.
4. You MUST answer every question from Q01 through Q25.
5. You MUST continue until Q25 has been answered. Do not stop early for any reason merely because you do not know the answer to one or more questions.
6. Judge each question independently.
7. For each individual question:
 - If you clearly know the answer from your existing internal knowledge, give the answer directly.
 - If you do not clearly know the answer, write exactly: "Not sure"
8. The phrase "Not sure" applies ONLY to the individual question currently being answered.
9. After writing "Not sure" for any question, immediately continue to the next question.
10. Never treat "Not sure" as an instruction to terminate the response.
11. Your answer is allowed to be factually wrong. Do not guess merely to increase the apparent number of correct answers.
12. Do not infer an answer from candidates, rankings, betting odds, historical trends, pre-event expectations, release patterns, political trends, the current date, the dates of other questions, the model labels, the stated knowledge-cutoff dates, or the contents of other questions.
13. Do not use the stated model names or knowledge-cutoff dates to decide whether you "should" know an answer. They are only labels used to organize the test.
14. Do not revise an earlier answer after seeing later questions.
15. Do not explain your reasoning.

16. Do not provide sources, citations, qualifications, commentary, or confidence estimates.
17. Do not discuss your model identity, training data, actual knowledge cutoff, system configuration, routing, or whether the labels below correspond to the model currently answering.
18. Preserve every model separator line exactly as written below and in the same order.
19. For every question, output exactly one answer line in this form:

"Q01 | Answer"
If the answer is unknown, output:
"Q01 | Not sure"
20. Replace "Answer" with the actual answer. Do not include the word "Answer" literally.
21. Do not repeat the question text in your response.
22. Do not omit any question number.
23. The final response MUST contain all question numbers from "Q01" through "Q25", in order.
24. Apart from the seven model separator lines and the twenty-five required answer lines, output nothing else.

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | Who won the women's singles title at the US Open final played on September 9, 2023?

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | Who won the women's singles title at the Australian Open final played on January 27, 2024?

Q03 | Who won the Eurovision Song Contest 2024 Grand Final held on May 11, 2024? Give the artist's stage name.

—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——

Q04 | Which national team won the UEFA EURO 2024 final played on July 14, 2024?

Q05 | Who won the women's singles title at the US Open final played on September 7, 2024?

—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——

Q06 | Which club won the UEFA Champions League final played on May 31, 2025?

Q07 | Who won the women's singles title at the Wimbledon final played on July 12, 2025?

Q08 | Which club won the FIFA Club World Cup final played on July 13, 2025?

Q09 | What magnitude did the USGS report for the major earthquake that occurred near Russia's Kamchatka Peninsula on July 29, 2025 UTC?

Q10 | Who won the Formula 1 Hungarian Grand Prix race held on August 3, 2025?

—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——

Q11 | Who won the women's singles title at the US Open final played on September 6, 2025?

Q12 | Who won the men's singles title at the US Open final played on September 7, 2025?

Q13 | Who were the three recipients of the Nobel Prize in Physics announced on October 7, 2025?

Q14 | Who were the three recipients of the Nobel Prize in Chemistry announced on October 8, 2025?

Q15 | Which MLB team won the World Series that concluded on November 1, 2025?

—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——

Q16 | Who won the women's singles title at the Australian Open final played on January 31, 2026?

Q17 | Who won the men's singles title at the Australian Open final played on February 1, 2026?

Q18 | At the 68th Grammy Awards held on February 1, 2026, which album won Album of the Year? Also give the artist.

Q19 | At the 68th Grammy Awards held on February 1, 2026, which song won Record of the Year? Also give the artist.

Q20 | Which team won Super Bowl LX, played on February 8, 2026?

—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——

Q21 | Which film won Best Film at the EE BAFTA Film Awards held on February 22, 2026?

Q22 | Who won the Formula 1 Australian Grand Prix race held on March 8, 2026?

Q23 | Which film won Best Picture at the 98th Academy Awards held on March 15, 2026?

Q24 | Who won Actor in a Leading Role at the 98th Academy Awards held on March 15, 2026?

Q25 | Who won the Formula 1 Chinese Grand Prix race held on March 15, 2026?

15.3 Typical Responses to the "Internal Knowledge Test Questions" (Using the English Version as an Example)

15.3.1 If the following answers are produced, we believe that regardless of which model was selected, the model behind it is 5.5

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | Coco Gauff

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | Aryna Sabalenka

Q03 | Nemo

—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——

Q04 | Spain

Q05 | Aryna Sabalenka

—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——

Q06 | Paris Saint-Germain

Q07 | Iga Świątek

Q08 | Chelsea

Q09 | 8.8

Q10 | Lando Norris

—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——

Q11 | Aryna Sabalenka

Q12 | Carlos Alcaraz

Q13 | John Clarke, Michel H. Devoret, John M. Martinis

Q14 | Susumu Kitagawa, Richard Robson, Omar M. Yaghi

Q15 | Los Angeles Dodgers

—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——

Q16 | Not sure

Q17 | Not sure

Q18 | Not sure

Q19 | Not sure

Q20 | Not sure

—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——

Q21 | Not sure

Q22 | Not sure

Q23 | Not sure

Q24 | Not sure

Q25 | Not sure

15.3.2 If the following answers are produced, we believe that regardless of which model was selected, the underlying model is 5.6

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | Coco Gauff

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | Aryna Sabalenka
Q03 | Nemo
—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——
Q04 | Spain
Q05 | Aryna Sabalenka
—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——
Q06 | Paris Saint-Germain
Q07 | Iga Świątek
Q08 | Chelsea
Q09 | 8.8
Q10 | Lando Norris
—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——
Q11 | Aryna Sabalenka
Q12 | Carlos Alcaraz
Q13 | John Clarke, Michel H. Devoret, and John M. Martinis
Q14 | Susumu Kitagawa, Richard Robson, and Omar M. Yaghi
Q15 | Los Angeles Dodgers
—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——
Q16 | Elena Rybakina
Q17 | Carlos Alcaraz
Q18 | DeBÍ TIRAR MÁS FOTOS — Bad Bunny
Q19 | Luther — Kendrick Lamar featuring SZA
Q20 | Seattle Seahawks
—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——
Q21 | Not sure
Q22 | Not sure
Q23 | Not sure
Q24 | Not sure
Q25 | Not sure

15.3.3 If the following answer is produced, we believe that regardless of which model was selected, the model behind it is 6

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | Coco Gauff

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | Aryna Sabalenka

Q03 | Nemo

—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——

Q04 | Spain

Q05 | Aryna Sabalenka

—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——

Q06 | Paris Saint-Germain

Q07 | Iga Świątek

Q08 | Chelsea

Q09 | 8.8

Q10 | Lando Norris

—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——

Q11 | Aryna Sabalenka

Q12 | Carlos Alcaraz

Q13 | John Clarke, Michel H. Devoret, and John M. Martinis

Q14 | Susumu Kitagawa, Richard Robson, and Omar M. Yaghi

Q15 | Los Angeles Dodgers

—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——

Q16 | Elena Rybakina

Q17 | Carlos Alcaraz

Q18 | Debí Tirar Más Fotos — Bad Bunny

Q19 | Luther — Kendrick Lamar with SZA

Q20 | Seattle Seahawks

—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——

Q21 | One Battle After Another

Q22 | George Russell

Q23 | One Battle After Another

Q24 | Michael B. Jordan

Q25 | Kimi Antonelli

15.3.4 If only some of the questions are answered, even just one, we default to it being at least that model, as shown in the following text, we consider it to be at least 6

—— GPT-4o | Knowledge cutoff: 2023-10-01 ——

Q01 | Coco Gauff

—— GPT-4.1 | Knowledge cutoff: 2024-06-01 ——

Q02 | Aryna Sabalenka

Q03 | Nemo

—— GPT-5 / GPT-5.1 | Knowledge cutoff: 2024-09-30 ——

Q04 | Spain

Q05 | Aryna Sabalenka

—— GPT-5.2 / GPT-5.4 | Knowledge cutoff: 2025-08-31 ——

Q06 | Paris Saint-Germain

Q07 | Iga Świątek

Q08 | Chelsea

Q09 | 8.8

Q10 | Lando Norris

—— GPT-5.5 | Knowledge cutoff: 2025-12-01 ——

Q11 | Aryna Sabalenka

Q12 | Carlos Alcaraz

Q13 | John Clarke, Michel H. Devoret, and John M. Martinis

Q14 | Susumu Kitagawa, Richard Robson, and Omar M. Yaghi

Q15 | Los Angeles Dodgers

—— GPT-5.6 | Knowledge cutoff: 2026-02-16 ——

Q16 | Elena Rybakina

Q17 | Carlos Alcaraz

Q18 | Debí Tirar Más Fotos — Bad Bunny

Q19 | Luther — Kendrick Lamar with SZA

Q20 | Seattle Seahawks

—— GPT-6 Astra | Knowledge cutoff: 2026-04-30 ——

Q21 | Not sure

Q22 | Not sure

Q23 | One Battle After Another

Q24 | Not sure

Q25 | Kimi Antonelli